

# Top 20+ RAG Interview Questions



Yana Khare

Last Updated : 22 Jan, 2025



RAG has been a game-changer in the developing fields of [Generative AI](#), Data Science, and AI. Because RAG models let machines produce more accurate, coherent, and consistent language with facts, they transform how humans engage with technology. RAG is bringing the idea of robots that can write unique content, engrossing product descriptions, and news pieces to life. Even with RAG's increasing importance, prospective data scientists and AI enthusiasts still need access to comprehensive information. This article fills that knowledge gap by offering the top 20+ RAG interview questions.

## Table of contents

1. RAG Interview Questions for Beginners
2. Intermediate Level RAG Interview Questions
3. Difficult Level RAG Interview Questions

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

Show details

Accept all cookies

Use necessary cookies

Free Course



# Building and Evaluating RAG System

Build a RAG-based Q&A app from scratch • Diagnose pain points • Tune retrieval metrics

[Enroll Now](#)

## RAG Interview Questions for Beginners

### Q1. What is Retrieval-Augmented Generation (RAG)?

A. [Retrieval-Augmented Generation \(RAG\)](#) is an approach that combines retrieval-based methods with generative models to enhance the performance of NLP tasks. In RAG, a retriever component first searches through a large corpus of documents to find relevant information based on the input query. Then, a generative model uses this retrieved information to generate a response or output. This two-step process allows RAG to leverage both the precision of retrieval methods and the flexibility of generative models. Therefore, it is particularly effective for tasks that require understanding and generating natural language based on external knowledge.

### Q2. Can you explain the basic difference between RAG and traditional language models?

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

external sources after training. This means they cannot answer questions or provide details about topics or events that were not part of their training data.

On the other hand, RAG (Retrieval-Augmented Generation) combines the strengths of traditional language models with a retrieval mechanism. It first retrieves relevant information from an external corpus of documents (e.g., databases, articles, or the web) and then uses that retrieved information to generate a response. This allows RAG to access up-to-date or domain-specific knowledge, making it more contextually aware and capable of providing accurate and informative responses, even for topics outside its original training data.

### Q3. What are some common applications of RAG in AI?

A. RAG (Retrieval-Augmented Generation) has various applications across different domains in AI, including:

1. **Question-Answering Systems:** RAG is widely used to build systems that provide accurate and contextually relevant answers to user queries. It retrieves relevant information from external sources (e.g., databases, documents, or the web) and generates precise responses, making it ideal for tasks like customer support, educational tools, and knowledge-based systems.
2. **Agentic RAG:** RAG can be integrated into agentic systems, where AI agents perform complex tasks by retrieving and synthesizing information from external sources. This allows them to make informed decisions, provide

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

3. **Content Generation:** RAG can generate high-quality, informative content by retrieving and combining information from multiple sources. It is particularly useful for tasks like summarizing articles, creating reports, or generating contextually rich narratives, ensuring the output is both coherent and factually accurate.
4. **Knowledge-Intensive Tasks:** RAG excels in domains requiring access to up-to-date or domain-specific knowledge, such as legal research, medical diagnosis, or technical support. By retrieving relevant documents or data, it ensures the generated responses are grounded in accurate and reliable information.

#### Q4. How does RAG improve the accuracy of responses in AI models?

A. RAG improves the accuracy of responses in AI models by leveraging a two-step approach that combines retrieval-based methods with generative models. The retrieval component first searches through a large corpus of documents to find relevant information based on the input query. This ensures that the system has access to up-to-date, domain-specific, or factual knowledge that may not be present in the model's training data. The generative component then uses this retrieved information to craft a coherent and contextually appropriate response. By incorporating external knowledge, RAG provides more accurate, informative, and reliable responses compared to traditional generative models that rely solely on learned patterns from their training data.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

responsible for searching the available data based on the input query and retrieving relevant documents. The retrieved documents then serve as the basis for the generative model to generate accurate and informative responses. The significance of retrievers lies in their ability to provide access to external knowledge, enhancing the context awareness and accuracy of RAG systems.

## Q6. What types of data sources are typically used in RAG systems?

A. In RAG systems, various types of data sources can be used, including:

- **Web Sources:** RAG systems can also retrieve information from the web by accessing online databases, websites, or search engine results. This allows the system to gather up-to-date or real-time information for generating responses.
- **Document Corpora:** RAG systems commonly use collections of text documents, such as books, research papers, articles, and websites. These corpora provide a rich source of unstructured or semi-structured information that the retriever can search through to find relevant content.
- **Knowledge Bases:** RAG systems can also utilize knowledge bases, such as wikis, encyclopedias, or domain-specific repositories. While these sources are not strictly structured databases, they often contain well-organized factual information that can be retrieved and used to generate accurate responses.

---

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

## Q7. How does RAG contribute to the field of conversational AI?

A. By allowing conversational agents to access and use outside knowledge sources, RAG advances [conversational AI by improving the agents](#)' capacity to produce insightful and contextually appropriate replies while interacting with others. By integrating generative models and retrieval-based techniques, RAG makes it possible for conversational agents to comprehend and react to user inquiries more precisely, resulting in more meaningful and captivating exchanges.

## Q8. What is the role of the retrieval component in RAG?

A. Based on the input question, the retrieval component of RAG searches through the available data sources, such as document corpora or knowledge bases, to identify pertinent information. This component finds and retrieves documents or data points containing relevant information using various retrieval approaches, including keyword matching and semantic search. The generative model receives and uses the relevant data retrieved to generate a response. The retrieval component dramatically increases RAG systems' accuracy and context awareness by making external knowledge more accessible.

## Intermediate Level RAG Interview Questions

### Q9. How does RAG handle bias and misinformation?

A. RAG can help mitigate bias and misinformation by leveraging a two-step

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

the retrieved information before generating a response. Thereby reducing biased or inaccurate information propagation. RAG aims to provide more reliable and accurate responses by incorporating external knowledge sources and validation mechanisms.

## Q10. What are the benefits of using RAG over other NLP techniques?

A. Some of the key benefits of using RAG over other NLP techniques include:

- **Enhanced Accuracy:** Utilizing external knowledge sources, RAG can produce more accurate and contextually appropriate replies than standard language models.
- **Context-Awareness:** RAG's retrieval component enables it to comprehend and consider a query's context, producing more meaningful and persuasive answers.
- **Flexibility:** RAG is a flexible solution for a broad range of [NLP applications](#). It can be tailored to different tasks and domains using multiple data sources.
- **Bias and Misinformation Mitigation:** RAG may help reduce bias and misinformation by prioritizing reliable sources and confirming retrieved information.

## Q11. Can you discuss a scenario where RAG would be particularly useful?

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

and reliable healthcare websites to get pertinent information. Afterward, the generative model would use this knowledge to provide replies relevant to the user's context and instructive.

RAG has the potential to enhance the precision and dependability of the healthcare chatbot by integrating external knowledge sources with generating capabilities. This would guarantee that users obtain reliable and current medical information. This approach can enhance the user experience, build trust with users, and provide valuable support in accessing reliable healthcare information.

## Q12. How does RAG integrate with existing machine learning pipelines?

A. Developers can integrate RAG into existing [machine learning pipelines](#) by using it as a component responsible for handling natural language processing tasks. Typically, they can connect the retrieval component of RAG to a database or document corpus, where it searches for relevant information based on the input query. Subsequently, the generative model processes the retrieved information to generate a response. This seamless integration allows RAG to leverage existing data sources and infrastructure, making it easier to incorporate into various machine learning pipelines and systems.

## Q13. What challenges does RAG solve in natural language processing?

A. RAG addresses several challenges in [natural language processing](#), including:

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details



- **Information Retrieval:** By leveraging retrieval-based methods, RAG can efficiently search through large datasets or document corpora to retrieve relevant information, improving the accuracy and relevance of generated responses.
- **Bias and Misinformation:** As discussed earlier, RAG can help mitigate bias and misinformation by prioritizing credible sources and validating retrieved information, enhancing the reliability of the generated content.
- **Personalization:** RAG can be adapted to personalize responses based on user preferences or historical interactions by retrieving and utilizing relevant information from previous interactions or user profiles.

#### Q14. How does RAG ensure the retrieved information is up-to-date?

A. Ensuring that retrieved information is up-to-date is crucial for the accuracy and reliability of RAG systems. To address this, developers can design RAG to regularly update its database or document corpus with the latest information from reputable and credible sources. They can also configure the retrieval component to prioritize recent publications or updates when searching for relevant information. Implementing continuous monitoring and updating mechanisms allows them to refresh the data sources and ensure the retrieved information remains current and relevant.

#### Q15. Can you explain how RAG systems are developed and used?

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

## 1. **Pre-trained Language Model (LLM):**

The generative component of RAG is usually a pre-trained large language model (LLM), such as GPT or similar transformer-based architectures. These models are pre-trained on massive text corpora to understand language patterns, structures, and representations. This pre-training enables the model to generate coherent and contextually relevant text.

## 2. **Retriever Component:**

The retriever in RAG is not trained but is instead an algorithm designed to search through a dataset or document corpus (e.g., Wikipedia, web pages, or domain-specific documents) to find relevant information based on the input query. Common retrieval algorithms include dense retrieval methods (e.g., using embeddings) or sparse retrieval methods (e.g., BM25).

## 3. **Integration of Components:**

RAG systems combine the retriever and the pre-trained LLM. When a query is received, the retriever first identifies relevant documents or passages from the corpus. The retrieved information is then passed to the LLM, which generates a response based on both the query and the retrieved context.

## 4. **Fine-tuning (Optional):**

In some cases, the LLM may be fine-tuned specifically for the RAG use case to better handle the integration of retrieved information. However, fine-tuning is not always necessary, as one of the key advantages of RAG is its ability to leverage pre-trained models and external knowledge without requiring costly

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

general capabilities and the retriever's ability to fetch relevant information, making RAG a cost-effective and flexible solution for knowledge-intensive tasks.

## Q16. What is the impact of RAG on the efficiency of language models?

A. RAG (Retrieval-Augmented Generation) significantly enhances the accuracy of language models by incorporating external knowledge through its retriever component, which fetches relevant information from large document corpora or datasets. This allows the generative model to produce more precise, contextually appropriate, and factually grounded responses, making RAG particularly effective for knowledge-intensive tasks. However, this improvement in accuracy comes with trade-offs in efficiency.

RAG systems typically have higher inference times compared to standalone generative models due to the additional step of retrieving relevant information, which adds computational overhead. While RAG avoids the need for costly fine-tuning of the generative model, it is more computationally expensive during inference than directly using a pre-trained generative model without retrieval. Despite this, RAG strikes a balance by reducing the generative model's reliance on memorized patterns from its training data, leading to more accurate responses. This makes RAG a powerful and scalable tool for real-world applications requiring high accuracy and contextual relevance, even with slightly higher computational costs.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

## Q17. How does RAG differ from fine-tuning?

**A.** RAG (Retrieval-Augmented Generation) and fine-tuning are two distinct approaches in natural language processing, each with its own strengths and use cases:

### **RAG (Retrieval-Augmented Generation)**

RAG enhances language models by combining retrieval-based methods with generative models. It uses a retriever component to fetch relevant information from external sources (e.g., document corpora or databases) and then leverages a generative model to produce responses based on the retrieved data. This approach allows RAG to access up-to-date or domain-specific knowledge without requiring modifications to the underlying generative model. RAG is particularly useful for tasks requiring external knowledge and avoids the need for costly fine-tuning.

### **Fine-Tuning:**

Fine-tuning involves adapting a pre-trained language model to a specific task or domain by further training it on a smaller, task-specific dataset. This process updates the model's parameters to improve its performance on the target task. While fine-tuning can significantly enhance a model's accuracy for specific applications, it is computationally expensive and requires labeled data. Additionally, fine-tuned models are limited to the knowledge present in their training data and cannot dynamically access external information.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

in their parameters from training.

- **Cost:** RAG avoids the computational expense of fine-tuning, making it more cost-effective for tasks requiring external knowledge. Fine-tuning, on the other hand, involves significant computational and data resources.
- **Flexibility:** RAG is more flexible for tasks requiring up-to-date or domain-specific information, whereas fine-tuned models are tailored to specific tasks but lack the ability to access external data dynamically.

## Q18. In what ways can RAG enhance human-AI collaboration?

A. RAG can enhance human-AI collaboration by:

- **Increasing Retrieval of Information:** RAG's retrieval component may access and retrieve pertinent material from big datasets or document corpora. Thus giving consumers thorough and precise answers to their inquiries.
- **Improving Context Understanding:** By keeping context consistent during a discussion, RAG may produce more meaningful and compelling replies. Therefore, interactions between humans and AI are made more seamless and meaningful.
- **Customizing Responses:** RAG may consider user choices and past interactions to provide customized answers that meet each person's requirements and preferences.

Overall RAG's ability to leverage external knowledge sources and generate

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

A. The technical architecture of a RAG system typically consists of two main components:

- **Retriever Component:** This component is responsible for searching through a dataset or document corpus to retrieve relevant information based on the input query. It uses retrieval techniques like keyword matching, semantic search, or neural retrievers to identify and extract pertinent data.
- **Generative Model:** After the data is obtained, it is sent to a generative model, such as a transformer-based architecture (like GPT), which uses the information to process it and respond. Based on the information gathered, this model is taught to comprehend and produce writing that resembles a person's.

Together, these two parts perform a two-step procedure. The [generative model employs](#) the relevant data the retriever has located and retrieved to provide an accurate and contextually relevant answer.

## Q20. How does RAG maintain context in a conversation?

A. RAG uses information acquired from past encounters or inside the present discussion to retain context in a discourse. The retriever component continuously searches for relevant data based on the ongoing conversation, ensuring the generative model has the context needed for coherent responses. This iterative process allows RAG to adapt to the evolving discussion, resulting in more engaging interactions.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy & Cookies Policy](#).

[Show details](#)

- **Computational Complexity:** The two-step process involving retrieval and generation can be computationally intensive. Hence, this leads to increased inference times and resource requirements.
- **Dependency on Data Quality:** RAG's performance relies heavily on the quality and relevance of the information retrieved. If the retriever component fails to retrieve accurate or pertinent data, it can impact the overall accuracy and reliability of the generated responses.
- **Scalability:** Managing and updating large document corpora or datasets can pose challenges in scalability and maintenance. This is especially true for real-time applications or systems with dynamic content.
- **Bias and Misinformation:** Like other AI models, RAG can inadvertently propagate biases present in the training data or retrieve and generate misinformation if not properly controlled or validated.

Despite these limitations, ongoing research in RAG seeks to overcome challenges and enhance its performance in various natural language processing tasks.

## Q22. How does RAG handle complex queries that require multi-hop reasoning?

A. By using its retrieval component to conduct iterative searches over several documents or data points to gradually obtain pertinent information, RAG may handle difficult questions that call for multi-hop reasoning. The retriever component may follow a logic path by getting data from one source. Further, it

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

answers to intricate questions involving multi-hop reasoning in addition to piecing together fragmented information from several sources.

### Q23. Can you discuss the role of knowledge graphs in RAG?

A. Knowledge graphs play a critical role in RAG. They facilitate more accurate and efficient information retrieval and reasoning by offering organized representations of knowledge and links between things. Knowledge graphs may be included in RAG's retriever component to improve search capabilities by using the graph structure to traverse and retrieve pertinent information more efficiently. Using knowledge graphs, RAG may record and use semantic links between ideas and things. Thus enabling more contextually rich and nuanced answers to user inquiries.

### Q24. What are the ethical considerations when implementing RAG systems?

A. [Implementing RAG systems](#) raises several ethical considerations, including:

- **Bias and Fairness:** It is crucial to ensure that RAG systems do not perpetuate or amplify biases in the training data or retrieved information. Implementing measures to detect and mitigate bias can promote fairness and equity in the generated responses.
- **Accountability and Transparency:** Helping users understand how RAG systems work fosters responsibility and trust. Clear documentation and

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)



- **Privacy and Data Security:** When accessing and retrieving information from external sources, preserving user privacy and guaranteeing data security is critical. Strong data protection measures and abiding by privacy laws and standards can protect user data and maintain trust.
- **Accuracy and Reliability:** To prevent the spread of incorrect or misleading information, it is crucial to guarantee the correctness and dependability of the obtained data and the replies created. Enforcing quality assurance procedures and validation processes can help preserve the RAG system's integrity.
- **User Consent and Control:** Respecting user preferences and offering options for information access and personalization enhance user autonomy and consent in RAG systems.

## Conclusion

RAG is a testament to AI's boundless potential to change our world. It can improve human experiences and push the limits of what machines can comprehend and produce in terms of natural language. It is more than simply a technological breakthrough. Knowing about RAG is a trip worth taking, whether you're getting ready for an AI interview or just interested in what the future holds for AI. It will open your eyes to new and creative possibilities in the exciting field of [artificial intelligence](#).

Unlock the secrets to mastering RAG!  Join our [GenAI Pinnacle Program](#)

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

A 23-year-old, pursuing her Master's in English, an avid reader, and a melophile. My all-time favorite quote is by Albus Dumbledore - "Happiness can be found even in the darkest of times if one remembers to turn on the light."

[Career](#)[Intermediate](#)[Interview Prep](#)[Interviews](#)[Listicle](#)[RAG](#)

## Free Courses



### Generative AI - A Way of Life

Explore Generative AI for beginners: create text and images, use top AI tools, learn practical skills, and ethics.



### Getting Started with Large Language Models

Master Large Language Models (LLMs) with this course, offering clear guidance in NLP and model

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)



## Building LLM Applications using Prompt Engineering

This free course guides you on building LLM apps, mastering prompt engineering, and developing chatbots with enterprise data.



## Improving Real World RAG Systems: Key Challenges & Practical Solutions

Explore practical solutions, advanced retrieval strategies, and agentic RAG systems to improve context, relevance, and accuracy in AI-driven applications.



## Microsoft Excel: Formulas & Functions

Master MS Excel for data analysis with key formulas, functions, and LookUp tools in this comprehensive course.

## RECOMMENDED ARTICLES

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

[Unveiling Retrieval Augmented Generation \(RAG\)|...](#)

[RAG's Innovative Approach to Unifying Ret...](#)

[Top 8 Applications of RAGs in Workplaces](#)

[12 RAG Pain Points and their Solutions](#)

[What is Retrieval-Augmented Generation \(RAG\)?](#)

[How to Become a RAG Specialist in 2025?](#)

[Top 6 Books on Retrieval Augmented Generation \(...\)](#)

[5 Days Roadmap to Learn RAG](#)

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

What are your thoughts?...

Submit reply



surabhi2187

Very good interview questions and answers in simple language..thanks



We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

- Reach a Global Audience
- Share Your Expertise with the World
- Build Your Brand & Audience
- Join a Thriving AI Community
- Level Up Your AI Game
- Expand Your Influence in Generative AI



## Flagship Programs

GenAI Pinnacle Program | GenAI Pinnacle Plus Program | AI/ML BlackBelt Program |  
Agentic AI Pioneer Program

## Free Courses

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

Coding in Windsurf | Model Deployment using FastAPI | Building Data Analyst AI Agent | Getting started with OpenAI o3-mini | Introduction to Transformers and Attention Mechanisms

## Popular Categories

AI Agents | Generative AI | Prompt Engineering | Generative AI Application | News | Technical Guides | AI Tools | Interview Preparation | Research Papers | Success Stories | Quiz | Use Cases | Listicles

## Generative AI Tools and Techniques

GANs | VAEs | Transformers | StyleGAN | Pix2Pix | Autoencoders | GPT | BERT | Word2Vec | LSTM | Attention Mechanisms | Diffusion Models | LLMs | SLMs | Encoder Decoder Models | Prompt Engineering | LangChain | LlamaIndex | RAG | Fine-tuning | LangChain AI Agent | Multimodal Models | RNNs | DCGAN | ProGAN | Text-to-Image Models | DDPM | Document Question Answering | Imagen | T5 (Text-to-Text Transfer Transformer) | Seq2seq Models | WaveNet | Attention Is All You Need (Transformer Architecture) | WindSurf | Cursor

## Popular GenAI Models

Llama 4 | Llama 3.1 | GPT 4.5 | GPT 4.1 | GPT 4o | o3-mini | Sora | DeepSeek R1 | DeepSeek V3 | Janus Pro | Veo 2 | Gemini 2.5 Pro | Gemini 2.0 | Gemma 3 | Claude Sonnet 3.7 | Claude 3.5 Sonnet | Phi 4 | Phi 3.5 | Mistral Small 3.1 | Mistral NeMo | Mistral-7b | Bedrock | Vertex AI | Qwen QwQ 32B | Qwen 2 | Qwen 2.5 VL | Qwen Chat | Grok 3

## AI Development Frameworks

n8n | LangChain | Agent SDK | A2A by Google | SmolAgents | LangGraph | CrewAI | Agno | LangFlow | AutoGen | LlamaIndex | Swarm | AutoGPT

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

[Exploration](#) | [Big Data](#) | [Common Machine Learning Algorithms](#) | [Machine Learning](#) | [Google Data Science Agent](#)

## Company

[About Us](#)

[Contact Us](#)

[Careers](#)

## Learn

[Free Courses](#)

[AI&ML Program](#)

[Pinnacle Plus Program](#)

[Agentic AI Program](#)

## Contribute

[Become an Author](#)

[Become a Speaker](#)

[Become a Mentor](#)

[Become an Instructor](#)

## Discover

[Blogs](#)

[Expert Sessions](#)

[Learning Paths](#)

[Comprehensive Guides](#)

## Engage

[Community](#)

[Hackathons](#)

[Events](#)

[Podcasts](#)

## Enterprise

[Our Offerings](#)

[Trainings](#)

[Data Culture](#)

[AI Newsletter](#)

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)



[Terms & conditions](#) • [Refund Policy](#) • [Privacy Policy](#) • [Cookies Policy](#) © Analytics  
Vidhya 2025. All rights reserved.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)