

Home > Blog > Artificial Intelligence

Top 30 LLM Interview Questions and Answers for 2025

This article provides a comprehensive guide to large language model (LLM) interview questions, covering fundamental concepts, intermediate and advanced techniques, and specific questions for prompt engineers.

Jul 18, 2024 · 15 min read



Stanislav Karzhev
Research Intern at UCL

TOPICS

Artificial Intelligence

[Large language models](#) (LLMs) have become increasingly important in artificial intelligence, with applications across various industries.

As the demand for professionals with LLM expertise grows, this article provides a comprehensive set of interview questions and answers, covering fundamental concepts, advanced techniques, and practical applications.

If you're preparing for a job interview or simply want to expand your knowledge, this article will be useful.

Basic LLM Interview Questions

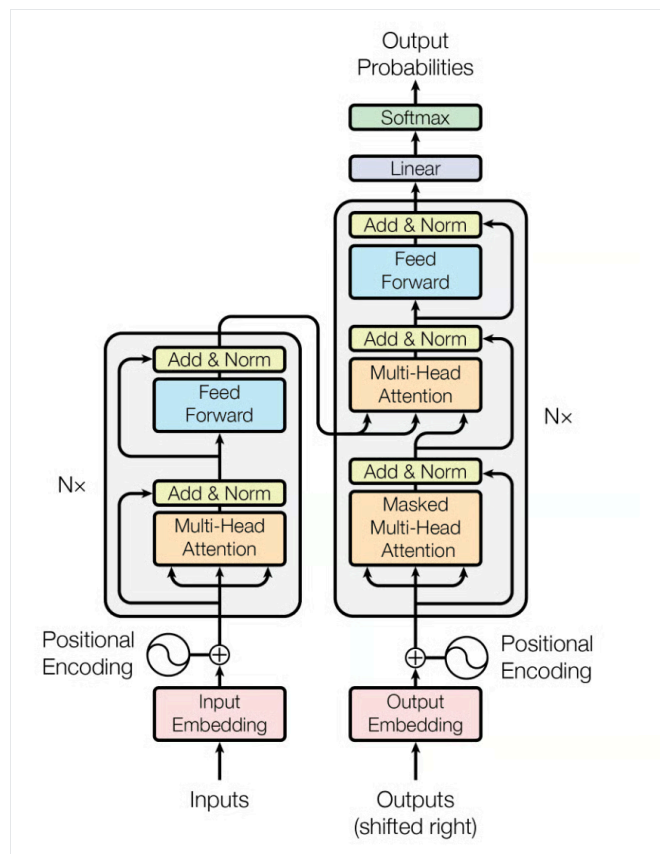
To understand LLMs, it's important to start with the fundamental concepts. These foundational questions cover essential aspects such as architecture, key mechanisms, and typical challenges, providing a solid base for learning more advanced topics.

What is the Transformer architecture, and how is it used in LLMs?

The [Transformer architecture](#) is a deep learning model introduced by [Vaswani et al.](#) in 2017, designed to handle sequential data with improved efficiency and performance than previous models like recurrent neural networks (RNNs) and long short-term memory (LSTMs).

It relies on [self-attention mechanisms](#) to process input data in parallel, making it highly scalable and capable of capturing long-range dependencies.

In LLMs, the Transformer architecture forms the backbone, enabling models to process large amounts of text data efficiently and generate contextually relevant and coherent text outputs.



The Transformer model architecture. [Source](#)

Explain the concept of "context window" in LLMs and its significance.

The context window in LLMs refers to the range of text (in terms of tokens or words) that the model can consider at once when generating or understanding language. The significance of the context window lies in its impact on the model's ability to generate logical and relevant responses.

A larger context window allows the model to consider more context, leading to better understanding and text generation, especially in complex or lengthy conversations. However, it also increases computational requirements, making it a balance between performance and efficiency.

What are some common pre-training objectives for LLMs, and how do they work?

Common pre-training objectives for LLMs include masked language modeling (MLM) and autoregressive language modeling. In MLM, random words in a sentence are masked, and the model is trained to predict the masked words based on the surrounding context. This helps the model understand the bidirectional context.

Autoregressive language modeling involves predicting the next word in a sequence and training the model to generate text one token at a time. Both objectives enable the model to learn language patterns and semantics from large corpora, providing a solid foundation for fine-tuning specific tasks.

Looking to get started with Generative AI?

Learn how to work with LLMs in Python right in your browser

[Start Now](#)



What is fine-tuning in the context of LLMs, and why is it important?

[Fine-tuning](#) in the context of LLMs involves taking a pre-trained model and further training it on a smaller, task-specific dataset. This process helps the model adapt its general language understanding to the nuances of the specific application, thereby improving performance.

This is an important technique because it leverages the broad language knowledge acquired during pre-training while modifying the model to perform well on specific applications, such as sentiment analysis, text summarization, or question-answering.

What are some common challenges associated with using LLMs?

Using LLMs comes with several challenges, including:

- **Computational resources:** LLMs require significant computational power and memory, making training and deployment resource-intensive.
- **Bias and fairness:** LLMs can inadvertently learn and propagate biases present in the training data, leading to unfair or biased outputs.
- **Interpretability:** Understanding and explaining the decisions made by LLMs can be difficult due to their complex and opaque nature.
- **Data privacy:** Using large datasets for training can raise concerns about data privacy and security.
- **Cost:** The development, training, and deployment of LLMs can be expensive, limiting their accessibility for smaller organizations.

How do LLMs handle out-of-vocabulary (OOV) words or tokens?

LLMs handle out-of-vocabulary (OOV) words or tokens using techniques like subword tokenization (e.g., Byte Pair Encoding or BPE, and WordPiece). These techniques break down unknown words into smaller, known subword units that the model can process.

This approach ensures that even if a word is not seen during training, the model can still understand and generate text based on its constituent parts, improving flexibility and robustness.

What are embedding layers, and why are they important in LLMs?

Embedding layers are a significant component in LLMs used to convert categorical data, such as words, into dense vector representations. These embeddings capture semantic relationships between words by representing them in a continuous vector space where similar words exhibit stronger proximity. The importance of embedding layers in LLMs includes:

- **Dimensionality reduction:** They reduce the dimensionality of the input data, making it more manageable for the model to process.
- **Semantic understanding:** Embeddings capture nuanced semantic meanings and relationships between words, enhancing the model's ability to understand and generate human-like text.
- **Transfer learning:** Pre-trained embeddings can be used across different models and tasks, providing a solid foundation of language understanding that can be fine-tuned for specific applications.

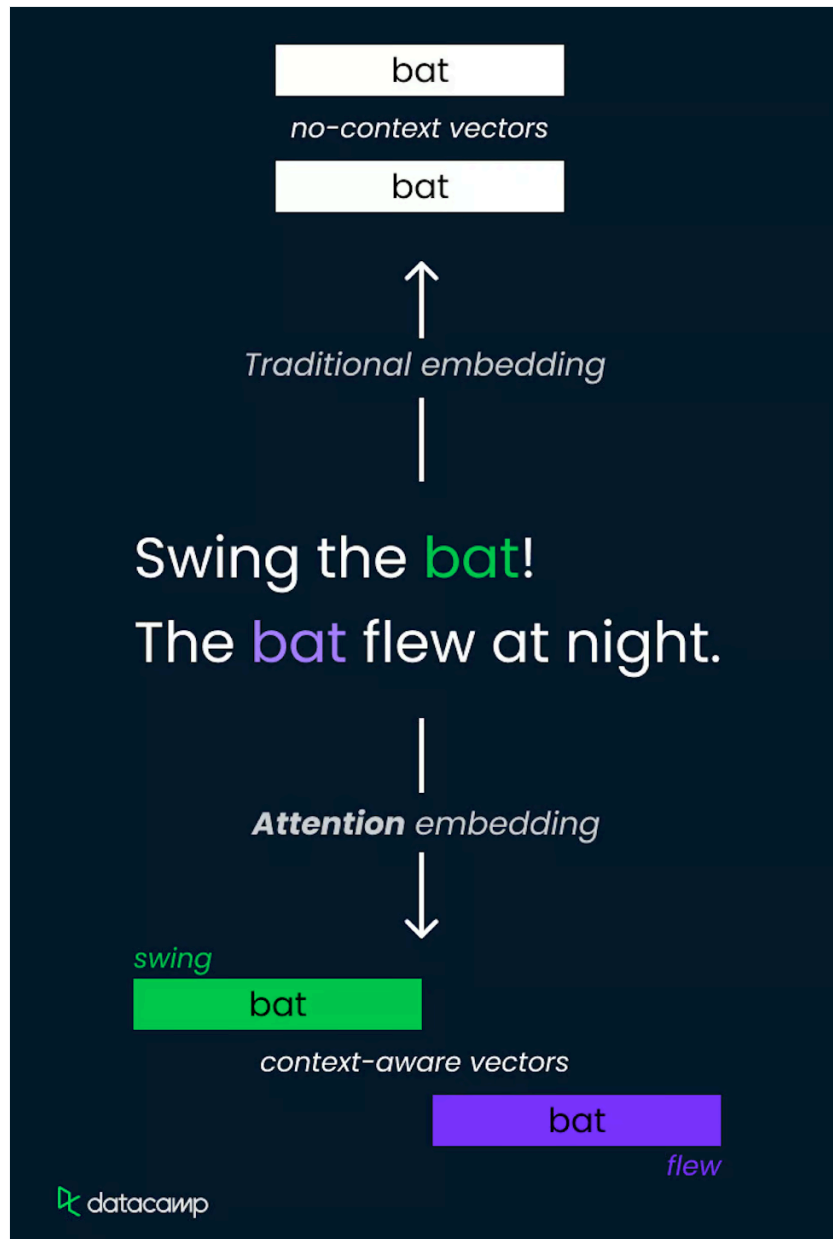
Intermediate LLM Interview Questions

Building upon basic concepts, intermediate-level questions delve into the practical techniques used to optimize LLM performance and address challenges related to computational efficiency and model interpretability.

Explain the concept of attention in LLMs and how it is implemented.

The concept of attention in LLMs is a method that allows the model to focus on different parts of the input sequence when making predictions. It dynamically assigns weights to other tokens in the input, highlighting the most relevant ones for the current task.

This is implemented using self-attention, where the model calculates attention scores for each token relative to all other tokens in the sequence, allowing it to capture dependencies regardless of their distance.



The self-attention mechanism is a core component of the Transformer architecture, enabling it to process information efficiently and capture long-range relationships.

What is the role of tokenization in LLM processing?

[Tokenization](#) converts raw text into smaller units called tokens, which can be words, subwords, or characters.

The role of tokenization in LLM processing is vital as it transforms text into a format that the model can understand and process.

Effective tokenization ensures that the model can handle a diverse range of inputs, including rare words and different languages, by breaking them down into manageable pieces. This step is necessary for optimal training and inference, as it standardizes the input and helps the model learn meaningful patterns in the data.

How do you measure the performance of an LLM?

Researchers and practitioners have developed numerous evaluation metrics to gauge the performance of an LLM. Common metrics include:

- **Perplexity:** Measures how well the model predicts a sample, commonly used in language modeling tasks.
- **Accuracy:** Used for tasks like text classification to measure the proportion of correct predictions.
- **F1 Score:** A harmonic mean of precision and recall, used for tasks like named entity recognition.
- **BLEU (Bilingual Evaluation Understudy) score:** Measures the quality of machine-generated text against reference translations, commonly used in machine translation.
- **ROUGE (Recall-Oriented Understudy for Gisting Evaluation):** A set of metrics that evaluate the overlap between generated text and reference text, often used in summarization tasks. They help quantify the model's effectiveness and guide further improvements.

What are some techniques for controlling the output of an LLM?

Several techniques can be used to control the output of an LLM, including:

- **Temperature:** Adjusting this parameter during sampling controls the randomness of the output. Lower temperatures produce more deterministic outputs, while higher values return more varied results.
- **Top-K sampling:** Limits the sampling pool to the top K most probable tokens, reducing the likelihood of generating less relevant or nonsensical text.
- **Top-P (nucleus) sampling:** Chooses tokens from the smallest set whose cumulative probability exceeds a threshold P, balancing diversity and coherence.
- **Prompt engineering:** Crafting specific prompts to guide the model towards generating desired outputs by providing context or examples.
- **Control tokens:** Using special tokens to signal the model to generate text in a specific style, format, or content type.

What are some approaches to reduce the computational cost of LLMs?

To reduce the computational cost of LLMs, we can employ:

- **Model pruning:** Removing less important weights or neurons from the model to reduce its size and computational requirements.
- **Quantization:** Converting the model weights from higher precision (e.g., 32-bit floating-point) to lower precision (e.g., 8-bit integer) reduces memory usage and speeds up inference.
- **Distillation:** Training a smaller model (student) to mimic the behavior of a larger, pre-trained model (teacher) to achieve similar performance with fewer resources.
- **Sparse attention:** Using techniques like sparse transformers to limit the attention mechanism to a subset of tokens, reduces computational load.
- **Efficient architectures:** Developing and using efficient model architectures specifically designed to minimize computational demands while maintaining performance, such as the Reformer or Longformer.

What is the importance of model interpretability in LLMs, and how can it be achieved?

Model interpretability is essential for understanding how an LLM makes decisions, which is important for building trust, ensuring accountability, and identifying and mitigating biases. Achieving interpretability can involve different approaches, such as:

- **Attention visualization:** Analyzing attention weights to see which parts of the input the model focuses on.
- **Saliency maps:** Highlighting input features that have the greatest influence on the model's output.

- **Model-Agnostic methods:** Using techniques like LIME (Local Interpretable Model-agnostic Explanations) to explain individual predictions.
- **Layer-wise relevance propagation:** Breaking down the model's predictions into contributions from each layer or neuron.

How do LLMs handle long-term dependencies in text?

LLMs handle long-term dependencies in text through their architecture, particularly the self-attention mechanism, which allows them to consider all tokens in the input sequence simultaneously. This ability to attend to distant tokens helps LLMs capture relationships and dependencies over long contexts.

Additionally, advanced models like the Transformer-XL and Longformer are specifically designed to extend the context window and manage longer sequences more effectively, ensuring better handling of long-term dependencies.

Advanced LLM Interview Questions

Understanding advanced concepts in LLMs is useful for professionals who aim to push the boundaries of what these models can achieve. This section explores complex topics and common challenges faced in the field.

Explain the concept of "few-shot learning" in LLMs and its advantages.

[Few-shot learning](#) in LLMs refers to the model's ability to learn and perform new tasks using only a few examples. This capability leverages the LLM's extensive pre-trained knowledge, enabling it to generalize from a small number of instances.

The primary advantages of few-shot learning include reduced data requirements, as the need for large task-specific datasets is minimized, increased flexibility, allowing the model to adapt to various tasks with minimal fine-tuning, and cost efficiency, as lower data requirements and reduced training times translate to significant cost savings in data collection and computational resources.

What are the differences between autoregressive and masked language models?

Autoregressive and masked language models differ mainly in their prediction approach and task suitability. Autoregressive models, like GPT-3 and GPT-4, predict the next word in a sequence based on the preceding words, generating text one token at a time.

These models are particularly well-suited for text-generation tasks. In contrast, masked language models, such as [BERT](#), randomly mask words in a sentence and train the model to predict these masked words based on the surrounding context. This bidirectional approach helps the model understand context from both directions, making it ideal for text classification and question-answering tasks.

How can you incorporate external knowledge into an LLM?

Incorporating external knowledge into an LLM can be achieved through several methods:

- [Knowledge graph integration](#): Augmenting the model's input with information from structured knowledge graphs to provide contextual information.
- [Retrieval-Augmented Generation \(RAG\)](#): Combines retrieval methods with generative models to fetch relevant information from external sources during text generation.
- [Fine-tuning with domain-specific data](#): Training the model on additional datasets that contain the required knowledge to specialize it for specific tasks or domains.
- **Prompt engineering**: Designing prompts that guide the model to utilize external knowledge effectively during inference.

What are some challenges associated with deploying LLMs in production?

Deploying LLMs in production involves various challenges:

- **Scalability:** Ensuring the model can handle large volumes of requests efficiently often requires significant computational resources and optimized infrastructure.
- **Latency:** Minimizing the response time to provide real-time or near-real-time outputs is critical for applications like chatbots and virtual assistants.
- **Monitoring and maintenance:** Continuously monitoring model performance and updating it to handle evolving data and tasks requires robust monitoring systems and regular updates.
- **Ethical and legal considerations:** Addressing issues related to bias, privacy, and compliance with regulations is essential to avoid ethical pitfalls and legal repercussions.
- **Resource management:** Managing the significant computational resources required for inference ensures cost-effectiveness and involves optimizing hardware and software configurations.

How do you handle model degradation over time in deployed LLMs?

Model degradation occurs when the performance of an LLM declines over time due to changes in the underlying data distribution. Handling model degradation involves regular retraining with updated data to maintain performance. Continuous monitoring is necessary to track the model's performance and detect signs of degradation.

Incremental learning techniques allow the model to learn from new data without forgetting previously learned information. Additionally, A/B testing compares the current model's performance with new versions and helps identify potential improvements before full deployment.

What are some techniques to ensure the ethical use of LLMs?

To ensure the [ethical use of LLMs](#), several techniques can be implemented:

- **Bias mitigation:** Applying strategies to identify and reduce biases in training data and model outputs, such as using balanced datasets and bias detection tools.
- **Transparency and explainability:** Developing models that provide interpretable and explainable outputs to foster trust and accountability, including using attention visualization and saliency maps.
- **User Consent and privacy:** Ensuring data used for training and inference complies with privacy regulations and obtaining user consent where necessary.
- **Fairness audits:** Conduct regular audits to evaluate the fairness and ethical implications of the model's behavior.
- **Responsible deployment:** Setting guidelines and policies for [responsible AI](#) deployment, including handling harmful or inappropriate content generated by the model.

How can you ensure the security of data used with LLMs?

Securing data used with LLMs requires implementing various measures. These include using encryption techniques for data at rest and in transit to protect against unauthorized access. Strict access controls are necessary to ensure that only authorized personnel can access sensitive data.

Anonymizing data to remove personally identifiable information (PII) before using it for training or inference is also crucial. Additionally, compliance with data protection regulations like GDPR or CCPA is essential to avoid legal issues.

These measures help protect data integrity, confidentiality, and availability. This protection is critical for maintaining user trust and adhering to regulatory standards.

Can you explain how techniques like reinforcement learning from human feedback (RLHF) can be used to improve the quality and safety of LLM

outputs, and what are some of the challenges associated with this approach?

RLHF is a technique that involves training an LLM to align its outputs with human preferences by incorporating feedback from human evaluators. This iterative process helps the model learn to generate responses that are not only accurate but also safe, unbiased, and helpful.

However, RLHF comes with challenges. One challenge is the potential for bias in the human feedback, as different evaluators might have varying preferences and interpretations.

Another challenge is the scalability of the feedback process, as collecting and incorporating large amounts of human feedback can be time-consuming and expensive. Additionally, ensuring that the reward model used in RLHF accurately captures the desired behaviors and values can be tricky.

Despite these challenges, RLHF has shown promising results in improving the quality and safety of LLM outputs, making it an important area of research and development in the field of prompt engineering.

More recently, an alternative to RLHF has emerged: [Reinforcement Learning From AI Feedback \(RLAIF\)](#).

LLM Interview Questions for Prompt Engineers

Prompt engineering is an important aspect of utilizing LLMs. It involves crafting precise and effective prompts to generate desired responses from the model. This section examines key questions that prompt engineers may encounter.

What is prompt engineering, and why is it crucial for working with LLMs?

Prompt engineering involves designing and refining prompts to guide LLMs in generating accurate and relevant outputs. It's vital for working with LLMs because the quality of the prompt directly impacts the model's performance.

Effective prompts can enhance the model's ability to understand the task, generate accurate and relevant responses, and reduce the likelihood of errors.

Prompt engineering is essential for maximizing the utility of LLMs in various applications, from text generation to complex problem-solving tasks.

Can you provide examples of different prompting techniques (zero-shot, few-shot, chain-of-thought) and explain when to use them?

- **Zero-shot prompting:** Provides the model with a task description without any examples. Typically used when there are no available examples or when we want to test the model's general understanding and flexibility.
- **Few-shot prompting:** Supplies a few examples along with the task description to guide the model. This is useful when the model needs context or examples to better understand the task.
- **Chain-of-Thought prompting:** Breaks down a complex task into smaller, sequential steps the model can follow. This can be beneficial for tasks that require logical reasoning and multi-step problem-solving.

How do you evaluate the effectiveness of a prompt?

Evaluating the effectiveness of a prompt involves:

- **Output quality:** Assessing the relevance, coherence, and accuracy of the model's responses.
- **Consistency:** Checking if the model consistently produces high-quality outputs across different inputs.
- **Task-specific metrics:** Using task-specific evaluation metrics, such as BLEU for translation or ROUGE for summarization, to measure performance.



- **Human evaluation:** Involve human reviewers to provide qualitative feedback on the performance.

What are some strategies for avoiding common pitfalls in prompt design (e.g., leading questions, ambiguous instructions)?

- **Avoid leading questions:** Ensure that prompts do not imply a specific answer, which can bias the model's response.
- **Clear and concise instructions:** Provide unambiguous and straightforward instructions to reduce confusion.
- **Context provision:** Include relevant context to help the model understand the task without overloading it with unnecessary information.
- **Iterative testing:** Continuously test and refine prompts based on the model's outputs and performance.

How do you approach iterative prompt refinement to improve LLM performance?

Iterative prompt refinement involves:

- **Initial design:** Start with a basic prompt based on task requirements.
- **Testing and evaluation:** Assess the prompt's performance using various metrics and obtain feedback.
- **Analysis:** Identify weaknesses or areas for improvement in the prompt.
- **Refinement:** Make adjustments to the prompt to address identified issues.
- **Repeat:** Repeat the testing and refinement process until the desired performance is achieved.

What tools or frameworks do you use to streamline the prompt engineering process?

Several tools and frameworks can streamline the prompt engineering process:

- **Interactive development environments (IDEs):** Using IDEs like Jupyter Notebook for easy experimentation and visualization.
- **APIs and SDKs:** Leveraging APIs and SDKs provided by LLM providers (e.g., OpenAI API) for prompt testing and deployment.
- **Automated testing frameworks:** Implementing automated testing frameworks to evaluate prompt performance across various scenarios and inputs.
- **Version control systems:** Using version control systems (e.g., Git) to track changes and collaborate on prompt engineering projects.
- **Visualization tools:** Employing visualization tools to analyze model outputs and identify patterns or areas for improvement.

How do you handle challenges like hallucination or bias in LLM outputs through prompt engineering?

This question addresses the ethical and practical issues of LLM-generated content. A strong answer would demonstrate awareness of these problems and discuss techniques like:

- **Fact verification prompts:** Incorporating prompts that encourage the model to verify its information against reliable sources.
- **Bias mitigation prompts:** Guiding the model to consider diverse perspectives and avoid discriminatory language.

- **Counterfactual prompts:** Asking the model to generate alternative scenarios or perspectives to challenge its initial assumptions.

Can you explain the role of prompt templates and how they are used in prompt engineering?

Prompt templates provide a structured format for prompts, often including placeholders for specific information or instructions. They can be reused across different tasks and scenarios, improving consistency and efficiency in prompt design.

A good answer would explain how prompt templates can be used to encapsulate best practices, incorporate domain-specific knowledge, and streamline the process of generating effective prompts for various applications.

How does the choice of a tokenizer impact prompt engineering and model performance?

The tokenizer plays a crucial role in how the LLM interprets and processes the input prompt. Different tokenizers have varying vocabulary sizes and handle out-of-vocabulary (OOV) words differently. A subword tokenizer like Byte Pair Encoding (BPE) can handle OOV words by breaking them into smaller subword units, while a word-based tokenizer might treat them as unknown tokens.

The choice of tokenizer can impact model performance in several ways. For instance, a subword tokenizer might be more effective in capturing the meaning of rare or technical terms, while a word-based tokenizer might be simpler and faster for general-purpose language tasks.

In prompt engineering, the choice of tokenizer can influence how you structure your prompts. For example, if you're using a subword tokenizer, you might need to pay more attention to how words are split into subwords to ensure that the model captures the intended meaning.

Earn a Top AI Certification

Demonstrate you can effectively and responsibly use AI.

Get Certified, Get Hired

Conclusion

This guide provided a set of interview questions to help you prepare for discussions on LLMs, ranging from basic principles to advanced strategies.

Whether you are preparing for an interview or looking to solidify your understanding, these insights will equip you with the knowledge needed to navigate and excel in the field of artificial intelligence.

If you want to read about the latest in AI and LLMs, I recommend these topics:

- [Claude Artifacts 101](#)
- [MatMul-Free LLMs: Key Concepts Explained](#)
- [SAMBA Hybrid Language Model: Key Concepts Explained](#)
- [What Is Buffer of Thoughts \(BoT\) and How It Works?](#)



AUTHOR

Stanislav Karzhev

in

Recent MSc graduate who specialises in Artificial Intelligence

TOPICS

Artificial Intelligence



Training more people?

Get your team access to the full DataCamp for business platform.

For BusinessFor a bespoke solution [book a demo](#).

Learn more about LLMs!

TRACK

AI Fundamentals

10hrs hr

Discover the fundamentals of AI, dive into models like ChatGPT, and decode generative AI secrets to navigate the dynamic AI landscape.

[See Details](#) →[Start Course](#)[See More](#) →

Related

BLOG

Top 30 AI Interview Questions and Answers For All Skill Levels



BLOG

Top 30 RAG Interview Questions and Answers for 2025



BLOG

Top 30 Generative AI Interview Questions and Answers for 2025

[See More](#) →

Grow your data skills with DataCamp for Mobile

Make progress on the go with our mobile courses and daily 5-minute coding challenges.



LEARN

[Learn Python](#)

[Learn AI](#)

[Learn Power BI](#)

[Learn Data Engineering](#)

[Assessments](#)

[Career Tracks](#)

[Skill Tracks](#)

[Courses](#)

[Data Science Roadmap](#)

DATA COURSES

[Python Courses](#)

[R Courses](#)

[SQL Courses](#)

[Power BI Courses](#)

[Tableau Courses](#)

[Alteryx Courses](#)

[Azure Courses](#)

[AWS Courses](#)

[Google Sheets Courses](#)

[Excel Courses](#)

[AI Courses](#)

[Data Analysis Courses](#)

[Data Visualization Courses](#)

[Machine Learning Courses](#)

[Data Engineering Courses](#)

[Probability & Statistics Courses](#)

DATALAB

[Get Started](#)

[Pricing](#)

[Security](#)

[Documentation](#)

CERTIFICATION

[Certifications](#)

[Data Scientist](#)

[Data Analyst](#)

[Data Engineer](#)

[SQL Associate](#)

[Power BI Data Analyst](#)

[Tableau Certified Data Analyst](#)

[Azure Fundamentals](#)

[AI Fundamentals](#)

RESOURCES

[Resource Center](#)

[Upcoming Events](#)

[Blog](#)

[Code-Alongs](#)

[Tutorials](#)

[Docs](#)

[Open Source](#)

[RDocumentation](#)

[Book a Demo with DataCamp for Business](#)

[Data Portfolio](#)

PLANS

[Pricing](#)

[For Students](#)

[For Business](#)

[For Universities](#)

[Discounts, Promos & Sales](#)

[Expense DataCamp](#)

[DataCamp Donates](#)

FOR BUSINESS

[Business Pricing](#)

[Teams Plan](#)

[Data & AI Unlimited Plan](#)

[Customer Stories](#)

[Partner Program](#)

ABOUT

[About Us](#)

[Learner Stories](#)

[Careers](#)

[Become an Instructor](#)

[Press](#)

[Leadership](#)

[Contact Us](#)

[DataCamp Español](#)

[DataCamp Português](#)

[DataCamp Deutsch](#)

[DataCamp Français](#)

SUPPORT

[Help Center](#)

[Become an Affiliate](#)



[Privacy Policy](#) [Cookie Notice](#) [Do Not Sell My Personal Information](#) [Accessibility](#) [Security](#) [Terms of Use](#)

© 2025 DataCamp, Inc. All Rights Reserved.