

Top 40 Machine Learning Interview Questions & Answers in 2025



Sakshi Raheja

Last Updated : 02 May, 2025



Ready to ace your machine learning interview? Start a journey through the intricacies of machine learning (ML) interview questions and answers, we delve into the fundamental concepts that underpin this dynamic field. From decoding the rationale behind F1 scores to navigating the nuances of logistic regression's nomenclature, these questions unveil the depth of understanding expected from ML enthusiasts. In this exploration, we unravel the significance of activation functions, the pivotal role of recall in cancer identification, and the impact of skewed data on model performance. Our quest spans diverse topics, from the principles of ensemble methods to the trade-offs inherent in the bias-variance interplay. As we unravel each question, the tapestry of ML knowledge unfolds, offering a holistic view of the intricate landscape of machine learning.

If you're a beginner, learn the basics of [machine learning here](#).

40 ML Interview Questions and Answers

Q1. Why do we take the harmonic mean of precision and recall when

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

Show details

Accept all cookies

Use necessary cookies

more than the arithmetic mean. This is crucial for cases where one of the metrics is significantly lower than the other. In classification tasks, precision and recall may have an inverse relationship; therefore, the harmonic mean ensures that the F1-score gives equal weight to precision and recall, providing a more balanced evaluation metric.

Q2. Why does Logistic regression have regression in its name even if it is used specifically for Classification?

A. Logistic regression doesn't directly classify but uses a linear model to estimate the probability of an event (0-1). We then choose a threshold (like 50%) to convert this to categories like 'yes' or 'no'. So, despite the 'regression' in its name, it ultimately tells us which class something belongs to.

Q3. What is the purpose of activation functions in neural networks?

A. [Activation functions](#) introduce non-linearity to neural networks, allowing them to learn complex patterns and relationships in data. Without activation functions, neural networks would reduce to linear models, limiting their ability to capture intricate features. Popular activation functions include sigmoid, tanh, and ReLU, each introducing non-linearity at different levels. These non-linear transformations enable neural networks to approximate complex functions, making them powerful tools for image recognition and natural language processing.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

A. In this scenario, Random Forest would be a more suitable choice. Logistic Regression is sensitive to the scale of input features, and unscaled features can affect its performance. On the other hand, Random Forest is less impacted by feature scaling due to its ensemble nature. Random Forest builds decision trees independently, and the scaling of features doesn't influence the splitting decisions across trees. Therefore, when dealing with unscaled data and limited insights, Random Forest would likely yield more reliable results.

Q5. In a binary classification problem aimed at identifying cancer in individuals, if you had to prioritize one performance metric over the other, considering you don't want to risk any person's life, which metric would you be more willing to compromise on, Precision or Recall, and why?

A. In identifying cancer, recall (sensitivity) is more critical than precision. Maximizing recall ensures that the model correctly identifies as many positive cases (cancer instances) as possible, reducing the chances of false negatives (missed cases). False negatives in cancer identification could have severe consequences. While precision is important to minimize false positives, prioritizing recall helps ensure a higher sensitivity to actual positive cases in the medical domain.

Q6. What is the significance of P-value when building a Machine Learning model?

A. P-values are used in traditional statistics to determine the significance of a

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

Q7. How does skewness in the distribution of a dataset affect the performance or behavior of machine learning models?

A. Skewness in the distribution of a dataset can significantly impact the performance and behavior of machine learning models. Here's an explanation of its effects and how to handle skewed data:

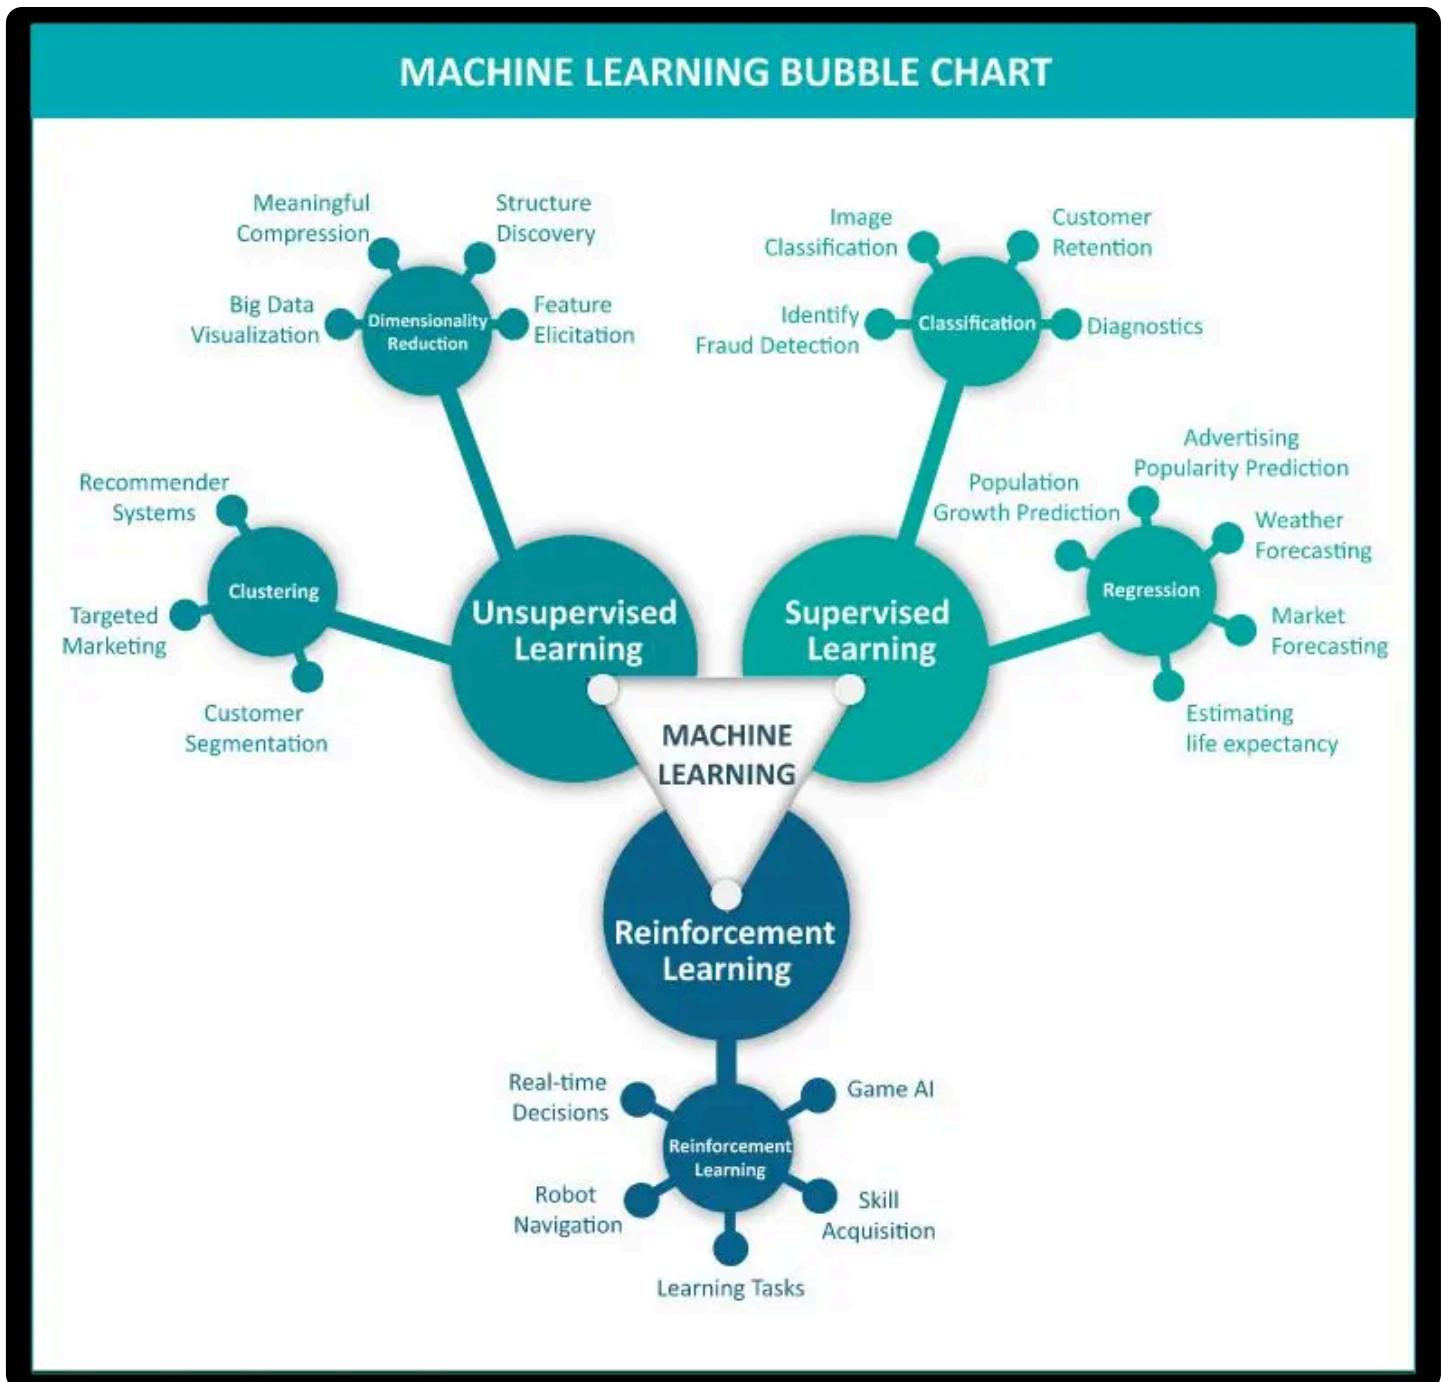
Effects of Skewed Data on Machine Learning Models:

- **Bias in Model Performance:** Skewed data can introduce bias in model training, especially with algorithms sensitive to class distribution. Models might be biased towards the majority class, leading to poor predictions for the minority class in classification tasks.
- **Impact on Algorithms:** Skewed data can affect the decision boundaries learned by models. For instance, in logistic regression or SVMs, the decision boundary might be biased towards the dominant class when one class dominates the other.
- **Prediction Errors:** Skewed data can result in inflated accuracy metrics. Models might achieve high accuracy by simply predicting the majority class yet fail to detect patterns in the minority class.

Also Read: [Machine Learning Algorithms](#)

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)



Q8. Describe a situation where ensemble methods could be useful.

A. Ensemble methods are particularly useful when dealing with complex and

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

weaknesses, an ensemble of models, such as Random Forest or Gradient Boosting, could be employed. Combining these models helps mitigate individual biases and uncertainties, resulting in a more reliable and accurate overall prediction.

Q9. How would you detect outliers in a dataset?

A. Outliers can be detected using various methods, including:

- **Z-Score:** Identify data points with a Z-score beyond a certain threshold.
- **IQR (Interquartile Range):** Flag data points outside the 1.5 times the IQR range.
- **Visualization:** Plotting box plots, histograms, or scatter plots can reveal data points significantly deviating from the norm.
- **Machine Learning Models:** Outliers may be detected using models trained to identify anomalies, like one-class SVMs or Isolation Forests.


Q10. Explain the Bias-Variance Tradeoff in Machine Learning. How does it impact model performance?

A. The bias-variance tradeoff refers to the delicate balance between the error introduced by bias and variance in machine learning models. A model with high bias oversimplifies the underlying patterns, leading to poor performance in training and unseen data. Conversely, a model with high variance captures noise

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

Support Vector Machines



Q11. Describe the working principle behind Support Vector Machines (SVMs) and their kernel trick. When would you choose SVMs over other algorithms?

A. SVMs aim to find the optimal hyperplane that separates classes with the maximum margin. The kernel trick allows SVMs to operate in a high-dimensional space, transforming non-linearly separable data into a linearly separable one.

Choose SVMs when:

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

- In scenarios where interpretability is less critical compared to predictive accuracy.

Q12. Explain the difference between lasso and ridge regularization.

A. Both lasso and ridge regularization are techniques to prevent overfitting by adding a penalty term to the loss function. The key difference lies in the type of penalty:

- **Lasso (L1 regularization):** Adds the absolute values of coefficients to the loss function, encouraging sparse feature selection. It tends to drive some coefficients to exactly zero.
- **Ridge (L2 regularization):** Adds the squared values of coefficients to the loss function. It discourages large coefficients but rarely leads to sparsity.

Choose lasso when feature selection is crucial and (overfitting) ridge when all features contribute meaningfully to the model.

Q13. Explain the concept of self-supervised learning in machine learning.

A. **Self-supervised learning** is a paradigm where models generate their labels from the existing data. It leverages the inherent structure or relationships within the data to create supervision signals without human-provided labels. Common self-supervised tasks include predicting missing parts of an image, filling in masked words in a sentence, or generating a relevant part of a video sequence.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

A. Bayesian optimization is an iterative model-based optimization technique that uses probabilistic models to guide the search for optimal hyperparameters. Unlike grid search or random search, Bayesian optimization considers the information gained from previous iterations, directing the search towards promising regions of the hyperparameter space. This approach is more efficient, requiring fewer evaluations, making it suitable for complex and computationally expensive models.

Q15. Explain the difference between semi-supervised and self-supervised learning.

- **Semi-Supervised Learning:** Involves training a model with both labeled and unlabeled data. The model learns from the labeled examples while leveraging the structure or relationships within the unlabeled data to improve generalization.
- **Self-Supervised Learning:** The model generates its labels from the existing data without external annotations. The learning task is designed so that the model predicts certain parts or features of the data, creating its supervision signals.

Q16. What is the significance of the out-of-bag error in machine learning algorithms?

A. The out-of-bag (OOB) error is a valuable metric in ensemble methods, particularly in Bagging (Bootstrap Aggregating). OOB error measures a model's

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

performance and can guide hyperparameter tuning for better predictive accuracy.

Q17. Explain the concept of Bagging and Boosting.

- **Bagging (Bootstrap Aggregating):** Bagging involves creating multiple subsets (bags) of the training dataset by randomly sampling with replacement. Each subset is used to train a base model independently. The final prediction aggregates predictions from all models, often reducing overfitting and improving generalization.
- **Boosting:** Boosting aims to improve the model sequentially by giving more weight to misclassified instances. It trains multiple weak learners, and each subsequent learner corrects the errors of its predecessors. Boosting, unlike bagging, is an adaptive method where each model focuses on the mistakes of the ensemble, leading to enhanced overall performance.

Also Read: [Ensemble Learning Methods](#)

Q18. What are the advantages of using Random Forest over a single decision tree?

- **Reduced Overfitting:** Random Forest mitigates overfitting by training multiple trees on different subsets of the data and averaging their predictions, providing a more generalized model.
- **Improved Accuracy:** The ensemble nature of Random Forest often results in

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

- **Feature Importance:** Random Forest measures feature importance, helping identify the most influential variables in the prediction process.
- **Robustness to Outliers:** Random Forest is less sensitive to outliers due to the averaging effect of multiple trees.

Q19. How does bagging reduce the variance of a model?

A. Bagging reduces model variance by training multiple instances of a base model on different subsets of the training data. The impact of individual outliers or noisy instances is diminished by averaging or combining the predictions of these diverse models. The ensemble's aggregated prediction tends to be more robust and less prone to overfitting specific patterns in a single subset of the data.

Q20. In bootstrapping and aggregating, can one sample from the data have one example (record) more than once? For example, can Row 344 of the dataset be included more than once in a single sample?

A. A sample can contain duplicates of the original data in bootstrapping. Since bootstrapping involves random sampling with replacement, some rows from the original dataset may be selected multiple times in a single sample. This characteristic contributes to the diversity of the base models in the ensemble.

Q21. Explain the connection between bagging and the “No Free Lunch” theorem in machine learning.

A. The “No Free Lunch” theorem states that no single machine learning algorithm

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

different subsets of data may require different models for optimal performance. Bagging provides a robust approach by leveraging the strengths of diverse models on different aspects of the data.

Q22. Explain the difference between hard and soft voting in a boosting algorithm.

- **Hard Voting:** In hard voting, each model in the ensemble makes a prediction, and the final prediction is determined by majority voting. The class with the most votes becomes the ensemble's prediction.
- **Soft Voting:** In soft voting, each model provides a probability estimate for each class, and the final prediction is based on the average or weighted average of these probabilities. Soft voting considers the confidence of each model's prediction.

Q23. How does voting boosting differ from simple majority voting and bagging?

- **Voting Boosting:** Boosting focuses on sequentially training weak learners, giving more weight to misclassified instances. Each subsequent model corrects errors, improving overall performance.
- **Simple Majority Voting:** In simple majority voting (as in bagging), each model has an equal vote, and the majority determines the final prediction. However, there's no sequential correction of errors.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

Q24. How does the choice of weak learners (e.g., decision stumps, decision trees) affect the performance of a voting-boosting model?

A. The choice of weak learners significantly impacts the performance of a voting-boosting model. Decision stumps (shallow trees with one split) are commonly used as weak learners. They are computationally less expensive and prone to underfitting, making them suitable for boosting. However, using more complex weak learners like deeper trees may lead to overfitting and degrade the model's generalization ability. The balance between simplicity and complexity in weak learners is crucial for boosting performance.

Q25. What is meant by forward and backward fill?

A. **Forward Fill:** Forward fill is a method used to fill missing values in a dataset by propagating the last observed non-missing value forward along the column. This method is useful when missing values occur intermittently in time-series or sequential data.

Backward Fill: Backward fill is the opposite, filling missing values by propagating the next observed non-missing value backward along the column. It is applicable in scenarios where future values are likely to be similar to past ones.

Both methods are commonly used in data preprocessing to handle missing values in time-dependent datasets.

Q26. Differentiate between feature selection and feature extraction.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

and efficiency. Methods include filter methods (based on statistical metrics), wrapper methods (using models to evaluate feature subsets), and embedded methods (incorporated into the model training process).

- **Feature Extraction:** [Feature extraction](#) transforms the original features into a new set of features, often of lower dimensionality. Techniques like Principal Component Analysis (PCA) and t-distributed Stochastic Neighbor Embedding (t-SNE) project data into a new space, capturing essential information while discarding less relevant details. Feature extraction is particularly useful when dealing with high-dimensional data or when feature interpretation is less critical.

Q27. How can cross-validation help in improving the performance of a model?

A. Cross-validation helps assess and improve model performance by evaluating how well a model generalizes to new data. It involves splitting the dataset into multiple subsets (folds), training the model on different folds, and validating it on the remaining folds. This process is repeated multiple times, and the average performance is computed. Cross-validation provides a more robust estimate of a model's performance, helps identify overfitting, and guides hyperparameter tuning for better generalization.

Q28. Differentiate between feature scaling and feature normalization. What are their primary goals and distinctions?

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

model training. Scaling methods include Min-Max Scaling, Z-score (standardization), and Robust Scaling.

- **Feature Normalization:** Feature normalization involves transforming features to a standard normal distribution with a mean of 0 and a standard deviation of 1 (Z-score normalization). It is a type of feature scaling that emphasizes achieving a specific distribution for the features.

Q29. Explain choosing an appropriate scaling/normalization method for a specific machine-learning task. What factors should be considered?

A. Choosing a scaling/normalization method depends on the characteristics of the data and the requirements of the machine-learning task:

- **Min-Max Scaling:** Suitable for algorithms sensitive to the scale of features (e.g., neural networks). Works well when data follows a uniform distribution.
- **Z-score Normalization (Standardization):** Suitable for algorithms assuming features are normally distributed. Resistant to outliers.
- **Robust Scaling:** Suitable when the dataset contains outliers. It scales features based on the interquartile range.

Consider the characteristics of the algorithm, the distribution of features, and the presence of outliers when selecting a method.

Q30. Compare and contrast z-scores with other standardization methods like min-max scaling.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

- **Min-Max Scaling:** Often, features are transformed to a specific range $[0, 1]$. Preserves the original distribution and is sensitive to outliers.

Both methods standardize features, but z-scores are suitable for normal distributions and robust to outliers. At the same time, min-max scaling is simple and applicable when preserving the original distribution is essential.

Q31. What is the IVF score, and what is its significance in building a machine-learning model?

A. "IVF score" is not a standard machine learning or feature engineering acronym. If "IVF score" refers to a specific metric or concept in a particular domain, additional context or clarification is needed to provide a relevant explanation.

Q32. How would you calculate the z-scores for a dataset with outliers? What additional considerations might be needed in such a case?

A. When calculating z-scores for a dataset containing outliers, it's crucial to be mindful of their influence on the mean and standard deviation, potentially skewing the z-score calculations. Outliers can significantly impact these statistics, leading to unreliable z-scores and misinterpretations of normality. To address this, one approach is to consider using robust measures such as the median absolute deviation (MAD) instead of the mean and standard deviation. MAD is less affected by outliers and provides a more resilient dispersion estimation. By employing MAD to compute the center and spread of the data, one can derive z-scores that are less susceptible to the influence of outliers, enabling

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

approach?

- **Pruning During Training:** During training, decision trees are grown to their full depth, and then unnecessary branches are pruned based on certain criteria (e.g., information gain). This helps prevent overfitting by removing branches that capture noise in the training data.
- **Pruning After Training:** The tree is allowed to grow without restrictions during training, and then pruning is applied afterward. This may involve removing nodes or branches that do not contribute significantly to overall predictive performance.

Advantages and Disadvantages:

- **Pruning During Training:** Pros include reduced overfitting and potentially more efficient training. However, it requires setting hyperparameters during training, which may lead to underfitting if not chosen appropriately.
- **Pruning After Training:** Allows the tree to capture more details during training and may improve accuracy. However, it may also lead to overfitting, and pruning decisions post-training might need to be more informed.

The choice depends on the dataset and the desired trade-off between model complexity and generalization.

Q34. Explain the core principles behind model quantization and pruning in machine learning. What are their main goals, and how do they differ?

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

to 8-bit integers. The primary goal is to reduce the model's memory footprint and computational requirements, making it more efficient for deployment on resource-constrained devices.

- **Pruning:** Model pruning involves removing unnecessary connections (weights) or entire neurons from a neural network. The main goal is to simplify the model structure, reduce the number of parameters, and improve inference speed. Pruning can be structured (removing entire neurons) or unstructured (removing individual weights).

Q35. How would you approach an Image segmentation problem?

A. Approaching an image segmentation problem involves the following steps:

- **Data Preparation:** Gather a labeled dataset with images and 32. How would you calculate the z-scores for a dataset with outliers? What additional considerations might be needed in such a case?
- **Robust Statistics:** Consider using robust statistics (e.g., median and interquartile range) instead of the mean and standard deviation to reduce the influence of outliers.
- **Outlier Treatment:** Evaluate whether to remove or transform outliers before calculating z-scores.corresponding pixel-level annotations indicating object boundaries.
- **Model Selection:** Choose a suitable segmentation model, such as U-Net,

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

- **Model Training:** Train the chosen model using the labeled dataset, optimizing for segmentation accuracy. Utilize pre-trained models if available for transfer learning.
- **Hyperparameter Tuning:** Fine-tune hyperparameters such as learning rate, batch size, and regularization to optimize model performance.
- **Evaluation:** Assess model performance using metrics like Intersection over Union (IoU) or Dice coefficient on a validation set.
- **Post-Processing:** Apply post-processing techniques to refine segmentation masks and handle potential artifacts or noise.

Q36. What is GridSearchCV?

A. [GridSearchCV](#), or Grid Search Cross-Validation, is a hyperparameter tuning technique in machine learning. It systematically searches through a predefined hyperparameter grid to find the combination that yields the best model performance. It performs cross-validation for each combination of hyperparameters, assessing the model's performance on different subsets of the training data.

The process involves defining a hyperparameter grid, specifying the machine learning algorithm, and selecting an evaluation metric. GridSearchCV exhaustively tests all possible hyperparameter combinations, helping identify the optimal set that maximizes model performance.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

- **False Positive (FP):** In binary classification, a false positive occurs when the model predicts the positive class incorrectly. It means the model incorrectly identifies an instance as belonging to the positive class when it belongs to the negative class.
- **False Negative (FN):** A false negative occurs when the model predicts the negative class incorrectly. It means the model fails to identify an instance that belongs to the positive class.

Significance:

- **False Positives:** In applications like medical diagnosis, a false positive can lead to unnecessary treatments or interventions, causing patient distress and additional costs.
- **False Negatives:** In critical scenarios like disease detection, a false negative may result in undetected issues, delaying necessary actions and potentially causing harm.

The significance depends on the specific context of the problem and the associated costs or consequences of misclassification.

Q38. What is PCA in Machine Learning, and can it be used for selecting features?

- **PCA (Principal Component Analysis):** PCA is a dimensionality reduction technique that transforms high-dimensional data into a lower-dimensional

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

informative components. However, there may be better choices for feature selection when the interpretability of individual features is crucial.

Q39. The model you have trained has a high bias and low variance. How would you deal with it?

Addressing a model with high bias and low variance involves:

- **Increase Model Complexity:** Choose a more complex model that can better capture the underlying patterns in the data. For example, move from a linear model to a non-linear one.
- **Feature Engineering:** Introduce additional relevant features the model may be missing to improve its learning ability.
- **Reduce Regularization:** If the model has regularization parameters, consider reducing them to allow it to fit the training data more closely.
- **Ensemble Methods:** Utilize ensemble methods, combining predictions from multiple models, to improve overall performance.
- **Hyperparameter Tuning:** Experiment with hyperparameter tuning to find the optimal settings for the model.

Q40. What is the interpretation of a ROC area under the curve?

A. The Receiver Operating Characteristic (ROC) curve is a graphical representation of a binary classification model's performance across different

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

- $AUC = 0.5$: The model performs no better than random chance.
- $AUC > 0.5$: The model performs better than random chance.

A higher AUC indicates better discrimination ability, with values closer to 1 representing superior performance. The ROC AUC is handy for evaluating models with class imbalance or considering different operating points.

Conclusion

In the tapestry of machine learning interview questions, we've traversed a spectrum of topics crucial for understanding the nuances of this evolving discipline. From the delicate balance of precision and recall in F1 scores to the strategic use of ensemble methods in diverse datasets, each question unraveled a layer of ML expertise. Whether discerning the criticality of recall in medical diagnoses or the impact of skewed data on model behavior, these questions probed the depth of knowledge and analytical thinking. As the journey concludes, it gives us a comprehensive understanding of ML's multifaceted landscape. It prepares us to navigate the challenges and opportunities that lie ahead in the dynamic realm of machine-learning interviews.



[Sakshi Raheja](#)

I am a passionate writer and avid reader who finds joy in weaving stories through the lens of data analytics and visualization. With a knack for blending creativity

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

A lover of both chai and coffee, I believe the right brew sparks creativity and sharpens focus—fueling my journey in the ever-evolving field of analytics. For me, every dataset holds a story, and I am always on a quest to uncover it.

[Intermediate](#)[Interview Prep](#)[Machine Learning](#)

Free Courses



Generative AI - A Way of Life

Explore Generative AI for beginners: create text and images, use top AI tools, learn practical skills, and ethics.



Getting Started with Large Language Models

Master Large Language Models (LLMs) with this course, offering clear guidance in NLP and model training made simple.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)



Building LLM Applications using Prompt Engineering

This free course guides you on building LLM apps, mastering prompt engineering, and developing chatbots with enterprise data.



Improving Real World RAG Systems: Key Challenges & Practical Solutions

Explore practical solutions, advanced retrieval strategies, and agentic RAG systems to improve context, relevance, and accuracy in AI-driven applications.



Microsoft Excel: Formulas & Functions

Master MS Excel for data analysis with key formulas, functions, and LookUp tools in this comprehensive course.

RECOMMENDED ARTICLES

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

[Top 15 Important Machine Learning Interview Que...](#)

[Top 20 Most Asked Machine Learning Interview Qu...](#)

[Top 40 Machine Learning Questions & Answers...](#)

[Get your Machine Learning Basics Right to Crack...](#)

[15 Most Common Data Science Interview Questions](#)

[Top 15 Important Data Science Interview Questions](#)

[8 Ways to Improve Accuracy of Machine Learning ...](#)

[15 Mlops Interview Questions for 2025](#)

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

[Show details](#)

What are your thoughts?...

Submit reply

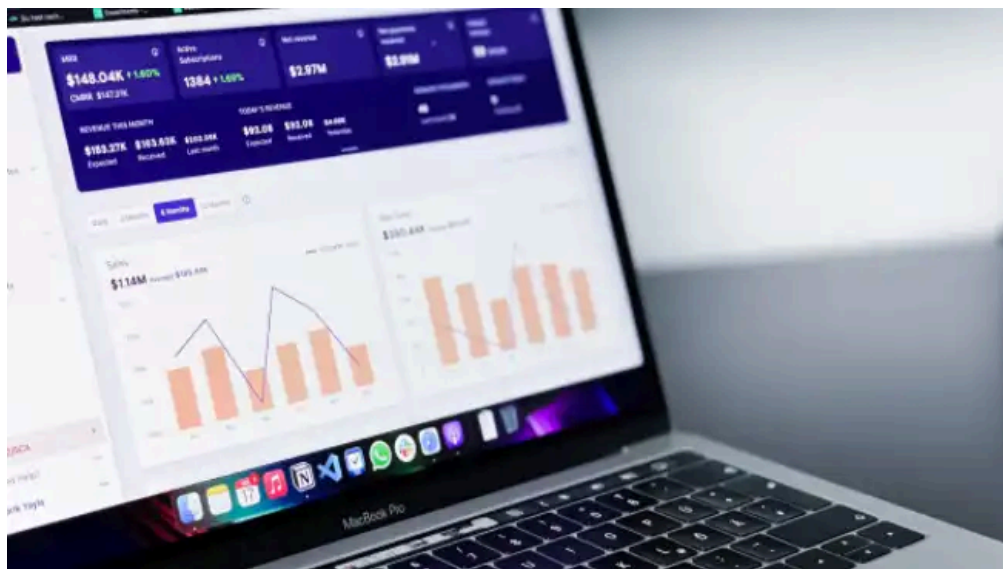
Become an Author →

Share insights, grow your voice, and inspire the data community.

- Reach a Global Audience
- Share Your Expertise with the World
- Build Your Brand & Audience
- Join a Thriving AI Community
- Level Up Your AI Game
- Expand Your Influence in Genrative AI

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details



We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)

Flagship Programs

GenAI Pinnacle Program | GenAI Pinnacle Plus Program | AI/ML BlackBelt Program | Agentic AI Pioneer Program

Free Courses

Generative AI | DeepSeek | OpenAI Agent SDK | LLM Applications using Prompt Engineering | DeepSeek from Scratch | Stability.AI | SSM & MAMBA | RAG Systems using LlamaIndex | Building LLMs for Code | Python | Microsoft Excel | Machine Learning | Deep Learning | Mastering Multimodal RAG | Introduction to Transformer Model | Bagging & Boosting | Loan Prediction | Time Series Forecasting | Tableau | Business Analytics | Vibe Coding in Windsurf | Model Deployment using FastAPI | Building Data Analyst AI Agent | Getting started with OpenAI o3-mini | Introduction to Transformers and Attention Mechanisms

Popular Categories

AI Agents | Generative AI | Prompt Engineering | Generative AI Application | News | Technical Guides | AI Tools | Interview Preparation | Research Papers | Success Stories | Quiz | Use Cases | Listicles

Generative AI Tools and Techniques

GANs | VAEs | Transformers | StyleGAN | Pix2Pix | Autoencoders | GPT | BERT | Word2Vec | LSTM | Attention Mechanisms | Diffusion Models | LLMs | SLMs | Encoder Decoder Models | Prompt Engineering | LangChain | LlamaIndex | RAG | Fine-tuning | LangChain AI Agent | Multimodal Models | RNNs | DCGAN | ProGAN | Text-to-Image Models | DDPM | Document Question Answering | Imagen | T5 (Text-to-Text Transfer Transformer) | Seq2seq Models | WaveNet | Attention Is All You Need (Transformer Architecture) | WindSurf | Cursor

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

AI Development Frameworks

n8n | LangChain | Agent SDK | A2A by Google | SmolAgents | LangGraph | CrewAI | Agno | LangFlow | AutoGen | LlamaIndex | Swarm | AutoGPT

Data Science Tools and Techniques

Python | R | SQL | Jupyter Notebooks | TensorFlow | Scikit-learn | PyTorch | Tableau | Apache Spark | Matplotlib | Seaborn | Pandas | Hadoop | Docker | Git | Keras | Apache Kafka | AWS | NLP | Random Forest | Computer Vision | Data Visualization | Data Exploration | Big Data | Common Machine Learning Algorithms | Machine Learning | Google Data Science Agent

Company

About Us
Contact Us
Careers

Discover

Blogs
Expert Sessions
Learning Paths
Comprehensive Guides

Learn

Free Courses
AI&ML Program

Engage

Community
Hackathons

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our Privacy Policy & Cookies Policy.

Show details

[Become an Author](#)

[Our Offerings](#)

[Become a Speaker](#)

[Trainings](#)

[Become a Mentor](#)

[Data Culture](#)

[Become an Instructor](#)

[AI Newsletter](#)

[Terms & conditions](#) • [Refund Policy](#) • [Privacy Policy](#) • [Cookies Policy](#) © Analytics
Vidhya 2025.All rights reserved.

We use cookies essential for this site to function well. Please click to help us improve its usefulness with additional cookies. Learn about our use of cookies in our [Privacy Policy](#) & [Cookies Policy](#).

[Show details](#)