

Top 50 MongoDB Interview Questions with Answers for 2025


Last Updated : 29 Jan, 2025

MongoDB is one of the most popular **NoSQL databases**, known for its **flexibility**, **scalability**, and ability to handle **unstructured data**. Many top companies rely on MongoDB, making it an essential skill for developers. Whether you're starting your career or looking to level up, preparing well for your **MongoDB interview** is the key to success.

In this article, we have compiled a complete list of MongoDB interview questions, covering everything from **basics** to **advanced concepts**. These questions will help you build **confidence**, understand key topics, and ace your interview. Start preparing now to secure your dream job!

Table of Content

- [MongoDB Basic Interview Questions](#)
- [MongoDB Intermediate Interview Questions](#)
- [MongoDB Query Based Interview Questions](#)



We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Got It !

MongoDB Basic Interview Questions

MongoDB Basic Interview Questions focus on **foundational knowledge**, helping you understand how MongoDB operates as a **NoSQL database**. These questions typically cover key concepts like the differences between SQL and **NoSQL databases**, how MongoDB structures and stores data, and basic operations like **CRUD** (Create, Read, Update, Delete).

1. What is MongoDB, and How Does It Differ from Traditional SQL Databases?

- **MongoDB** is a [NoSQL](#) database which means it does not use the **traditional table-based** relational database structure. Instead of it uses a flexible and document-oriented data model that stores data in **BSON** (Binary JSON) format.
- Unlike SQL databases that use rows and columns, MongoDB stores data as [JSON](#)-like documents, making it easier to handle unstructured data and providing greater flexibility in terms of schema design.

2. Explain BSON and Its Significance in MongoDB.

[BSON](#) (Binary JSON) is a **binary-encoded serialization** format used by MongoDB to store documents. BSON extends JSON by adding support for data types such as dates and binary data and it is designed to be efficient in both storage space and scan speed. The binary format allows MongoDB to be more efficient with data retrieval and storage compared to text-based JSON.

3. Describe the Structure of a MongoDB Document.

A MongoDB document is a set of **key-value** pairs similar to a JSON object. Each key is a string and the value can be a variety of data types including

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
{
  "_id": ObjectId("507f1f77bcf86cd799439011"),
  "name": "Alice",
  "age": 25,
  "address": {
    "street": "123 Main St",
    "city": "Anytown",
    "state": "CA"
  },
  "hobbies": ["reading", "cycling"]
}
```

4. What are Collections And Databases In MongoDB?

- **Database:** A container for collections, equivalent to a database in SQL.
- **Collection:** A group of documents, similar to tables in SQL, but schema-less.
- For example, a users collection can be part of the mydatabase database.

5. How Does MongoDB Ensure High Availability and Scalability?

- MongoDB ensures high **availability** and **scalability** through its features like **replica sets** and [sharding](#).
- Replica sets provide redundancy and failover capabilities by ensuring that data is always available.
- Sharding distributes data across multiple servers, enabling horizontal scalability to handle large volumes of data and high traffic loads.

6. Explain the Concept of Replica Sets in MongoDB.

- A **replica set** in MongoDB is a group of mongod instances that maintain the same data set.

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

- If the primary node fails, an automatic election process selects a new primary to maintain high availability.

7. What are the Advantages of Using MongoDB Over Other Databases?

- **Flexibility:** MongoDB's document-oriented model allows for dynamic schemas.
- **Scalability:** Built-in sharding enables horizontal scaling.
- **High Availability:** Replica sets provide redundancy and automatic failover.
- **Performance:** Efficient storage and retrieval with BSON format.
- **Ease of Use:** JSON-like documents make it easier for developers to interact with data.

8. How to Create a New Database and Collection in MongoDB?

To create a new database and collection in MongoDB, you can use the mongo shell:

```
use mydatabase
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

This command switches to **mydatabase** (creating it if it doesn't exist) and creates a new collection named **mycollection**.

9. What is Sharding, and How Does It Work in MongoDB?

Sharding is a method for distributing data across multiple servers in MongoDB. It allows for horizontal scaling by splitting large datasets into smaller, more manageable pieces called **shards**.

- Each shard is a separate database that holds a portion of the data.
- MongoDB automatically balances data and load across shards, ensuring efficient data distribution and high performance.

10. Explain the Basic Syntax of MongoDB CRUD Operations.

CRUD operations in MongoDB are used to create, read, update, and delete documents.

- **Create:** `db.collection.insertOne({ name: "Alice", age: 25 })`
- **Read:** `db.collection.find({ name: "Alice" })`
- **Update:** `db.collection.updateOne({ name: "Alice" }, { $set: { age: 26 } })`
- **Delete:** `db.collection.deleteOne({ name: "Alice" })`

11. How to Perform Basic Querying in MongoDB?

Basic querying in MongoDB involves using the `find` method to retrieve documents that match certain criteria.

Example:

```
db.collection.find({ age: { $gte: 20 } })
```

This query retrieves all documents from the collection where the **age** field is greater than or equal to 20.

An **index** in MongoDB is a data structure that improves the speed of data retrieval operations on a collection. You can create an index using the `createIndex` method.

For example, to create an index on the name field:

```
db.collection.createIndex({ name: 1 })
```

13. How Does MongoDB Handle Data Consistency?

MongoDB provides several mechanisms to ensure data consistency:

- **Journaling:** MongoDB uses write-ahead logging to maintain data integrity.
- **Write Concerns:** It specifies the level of acknowledgment requested from MongoDB for write operations (e.g., acknowledgment from primary only, or acknowledgment from primary and secondaries).
- **Replica Sets:** Replication ensures data is consistent across multiple nodes, and read concerns can be configured to ensure data consistency for read operations.

14. How to Perform Data Import and Export in MongoDB?

To perform data import and export in MongoDB, you can use the `mongoimport` and `mongoexport` tools. These tools allow you to import data from **JSON, CSV or TSV** files into MongoDB and export data from MongoDB collections to JSON or CSV files.

Import Data:

```
mongoimport --db mydatabase --collection mycollection --file data.json
```

This command imports data from `data.json` into the `mycollection` collection

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
mongoexport --db mydatabase --collection mycollection --out  
data.json
```

This command exports data from the mycollection collection in the mydatabase database to data.json.

15. What are MongoDB Aggregation Pipelines and How are They Used?

The [aggregation pipeline](#) is a framework for data aggregation, modeled on the concept of data processing pipelines. Documents enter a multi-stage pipeline that transforms the documents into aggregated results. Each stage performs an operation on the input documents and passes the results to the next stage.

```
db.orders.aggregate([  
  { $match: { status: "A" } },           // Stage 1: Filter  
  documents by status  
  { $group: { _id: "$cust_id", total: { $sum: "$amount" } } },  
  // Stage 2: Group by customer ID and sum the amount  
  { $sort: { total: -1 } }              // Stage 3: Sort by  
  total in descending order  
])
```

In this example:

- Stage 1 (\$match) filters documents by status "A".
- Stage 2 (\$group) groups documents by customer ID and calculates the total amount for each group.
- Stage 3 (\$sort) sorts the results by total amount in descending order.

Aggregation pipelines are powerful and flexible, enabling complex data processing tasks to be executed within MongoDB.

MongoDB Intermediate Interview Questions explore **advanced concepts** and **features**, such as **schema design**, **aggregation pipelines**, **indexing strategies**, and **transaction management**. These questions help gauge your ability to utilize MongoDB efficiently in more complex scenarios.

1. Describe the Aggregation Framework in MongoDB

- The Aggregation Framework in MongoDB is a powerful tool for performing data processing and transformation on documents within a collection.
- It works by passing documents through a multi-stage pipeline, where each stage performs a specific operation on the data, such as **filtering**, **grouping**, **sorting**, **reshaping** and **computing aggregations**.
- This framework is particularly useful for creating complex data transformations and analytics directly within the [database](#).

2. How to Perform Aggregation Operations Using MongoDB?

Aggregation operations in MongoDB are performed using the `aggregate` method. This method takes an array of pipeline stages, each stage representing a step in the data processing pipeline.

Example: Calculate total sales for each product:

```
db.sales.aggregate([
  { $match: { status: "completed" } }, // Filter completed
  sales
  { $group: { _id: "$product", totalSales: { $sum: "$amount" } } }, // Group by product and sum the sales amount
  { $sort: { totalSales: -1 } } // Sort by total sales in
  descending order
])
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Write Concern in MongoDB refers to the level of acknowledgment requested from MongoDB for write operations. It determines how many nodes must confirm the write operation before it is considered successful. Write concern levels range from "**acknowledged**" (default) to "**unacknowledged**," "**journal**," and various "**replica acknowledged**" levels.

The importance of write concern lies in balancing between data durability and performance. Higher write concern ensures data is safely written to disk and replicated, but it may impact performance due to the added latency.

4. What are TTL Indexes, and How are They Used in MongoDB?

TTL (Time To Live) Indexes in MongoDB are special indexes that automatically remove documents from a collection after a certain period. They are commonly used for data that needs to expire after a specific time, such as **session information**, **logs**, or **temporary data**. To create a TTL index, you can specify the expiration time in seconds

Example: Remove documents 1 hour after createdAt:

```
db.sessions.createIndex({ "createdAt": 1 }, {  
  expireAfterSeconds: 3600 })
```

This index will remove documents from the sessions collection 1 hour (3600 seconds) after the createdAt field's value.

5. How to Handle Schema Design and Data Modeling in MongoDB?

Schema design and **data modeling** in MongoDB involve defining how data is organized and stored in a document-oriented database. Unlike SQL databases, MongoDB offers flexible schema design, which can be both an advantage and a challenge. Key considerations for schema design include:

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

- **Document Structure:** Designing documents that align with application query patterns for efficient read and write operations.
- **Indexing:** Creating indexes to support query performance.
- **Data Duplication:** Accepting some level of data duplication to optimize for read performance.
- **Sharding:** Designing the schema to support sharding if horizontal scaling is required.

6. What is GridFS, and When is it Used in MongoDB?

GridFS is a specification for storing and retrieving large files in MongoDB. It is used when files exceed the BSON-document size limit of 16 MB or when you need to perform efficient retrieval of specific file sections.

GridFS splits a large file into smaller chunks and stores each chunk as a separate document within two collections: **fs.files** and **fs.chunks**. This allows for efficient storage and retrieval of large files, such as images, videos, or large datasets.

7. Explain the Differences Between WiredTiger and MMAPv1 Storage Engines

Feature	WiredTiger	MMAPv1
Concurrency	Document-level concurrency, allowing multiple operations simultaneously.	Collection-level concurrency, limiting performance under heavy write operations.
Compression	Supports data compression, reducing storage requirements.	Does not support data compression.

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Feature	WiredTiger	MMAPv1
	for most workloads.	workloads.
Journaling	Uses write-ahead logging for better data integrity.	Basic journaling; less advanced than WiredTiger.
Status	Modern and default storage engine.	Legacy engine, deprecated in favor of WiredTiger.
Implementation	Advanced implementation with additional features.	Simple implementation but lacks advanced features.

8. How to Handle Transactions in MongoDB?

MongoDB supports multi-document [ACID transactions](#) by allowing us to perform a series of read and write operations across multiple documents and collections in a transaction. This ensures data consistency and integrity. To use transactions we typically start a session, begin a transaction, perform the operations and then **commit** or **abort** the transaction.

Example in JavaScript:

```
const session = client.startSession();

session.startTransaction();

try {
  db.collection1.insertOne({ name: "Alice" }, { session });
  db.collection2.insertOne({ name: "Bob" }, { session });
  session.commitTransaction();
} catch (error) {
  session.abortTransaction();
}
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
    session.endSession();  
}
```

9. Describe the MongoDB Compass Tool and Its Functionalities

MongoDB Compass is a [graphical user interface](#) (GUI) tool for MongoDB that provides an easy way to visualize, explore, and manipulate your data. It offers features such as:

- **Schema Visualization:** View and analyze your data schema, including field types and distributions.
- **Query Building:** Build and execute queries using a visual interface.
- **Aggregation Pipeline:** Construct and run aggregation pipelines.
- **Index Management:** Create and manage indexes to optimize query performance.
- **Performance Monitoring:** Monitor database performance, including slow queries and resource utilization.
- **Data Validation:** Define and enforce schema validation rules to ensure data integrity.
- **Data Import/Export:** Easily import and export data between MongoDB and JSON/CSV files.

10. What is MongoDB Atlas, and How Does it Differ From Self-Hosted MongoDB?

MongoDB Atlas is a fully managed cloud database service provided by MongoDB. It offers automated deployment, scaling, and management of MongoDB clusters across various cloud providers ([AWS](#), [Azure](#), [Google Cloud](#)). Key differences from self-hosted MongoDB include:

- **Managed Service:** Atlas handles infrastructure management, backups, monitoring, and upgrades.

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

- **Global Distribution:** Deploy clusters across multiple regions for low-latency access and high availability.
- **Integrations:** Seamless integration with other cloud services and MongoDB tools.

11. How to Implement Access Control and User Authentication in MongoDB?

Access control and user authentication in MongoDB are implemented through a role-based access control (RBAC) system. You create users and assign roles that define their permissions. To set up access control:

- **Enable Authentication:** Configure MongoDB to require authentication by starting the server with `--auth` or setting `security.authorization` to `enabled` in the configuration file.
- **Create Users:** Use the `db.createUser` method to create users with specific roles.

```
db.createUser({
  user: "admin",
  pwd: "password",
  roles: [{ role: "userAdminAnyDatabase", db: "admin" }]
});
```

- **Assign Roles:** Assign roles to users that define their permissions, such as read, write, or admin roles for specific databases or collections.

12. What are Capped Collections, and When are They Useful?

Capped collections in MongoDB are fixed-size collections that automatically overwrite the oldest documents when the specified size limit is reached. They maintain insertion order and are useful for scenarios

where you need to store a fixed amount of recent data, such as logging.

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
db.createCollection("logs", { capped: true, size: 100000 });
```

13. Explain the Concept of Geospatial Indexes in MongoDB

Geospatial indexes in MongoDB are special indexes that support querying of geospatial data, such as locations and coordinates. They enable efficient queries for proximity, intersections, and other spatial relationships. MongoDB supports two types of geospatial indexes: **2d** for **flat** geometries and **2dsphere** for spherical geometries.

Example of creating a 2dsphere index:

```
db.places.createIndex({ location: "2dsphere" });
```

14. How to Handle Backups and Disaster Recovery in MongoDB?

Handling backups and disaster recovery in MongoDB involves regularly creating backups of your data and having a plan for restoring data in case of failure. Methods include:

- **Mongodump/Mongorestore:** Use the mongodump and mongorestore utilities to create and restore binary backups.
- **File System Snapshots:** Use file system snapshots to take consistent backups of the data files.
- **Cloud Backups:** If using MongoDB Atlas, leverage automated backups provided by the service.
- **Replica Sets:** Use replica sets to ensure data redundancy and high availability. Regularly test the failover and recovery process.

15. Describe the Process of Upgrading MongoDB to a Newer Version

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

- **Check Compatibility:** Review the release notes and compatibility changes for the new version.
- **Backup Data:** Create a backup of your data to prevent data loss.
- **Upgrade Drivers:** Ensure that your application drivers are compatible with the new MongoDB version.
- **Upgrade MongoDB:** Follow the official MongoDB upgrade instructions, which typically involve stopping the server, installing the new version, and restarting the server.
- **Test Application:** Thoroughly test your application with the new MongoDB version to identify any issues.
- **Monitor:** Monitor the database performance and logs to ensure a successful upgrade.

16. What are Change Streams in MongoDB, and How are They Used?

Change Streams in MongoDB allow applications to listen for real-time changes to data in collections, databases, or entire clusters. They provide a powerful way to implement event-driven architectures by capturing insert, update, replace, and delete operations. To use Change Streams, you typically open a change stream cursor and process the change events as they occur.

Example:

```
const changeStream = db.collection('orders').watch();
changeStream.on('change', (change) => {
  console.log(change);
});
```

This example listens for changes in the orders collection and logs the change events.

Hashed Sharding Keys in MongoDB distribute data across shards using a hashed value of the shard key field. This approach ensures an even distribution of data and avoids issues related to data locality or uneven data distribution that can occur with range-based sharding. Hashed sharding is useful for fields with monotonically increasing values, such as timestamps or identifiers.

Example:

```
db.collection.createIndex({ _id: "hashed" });  
sh.shardCollection("mydb.mycollection", { _id: "hashed" });
```

18. How to Optimize MongoDB Queries for Performance?

Optimizing MongoDB queries involves several strategies:

- **Indexes:** Create appropriate indexes to support query patterns.
- **Query Projections:** Use projections to return only necessary fields.
- **Index Hinting:** Use index hints to force the query optimizer to use a specific index.
- **Query Analysis:** Use the `explain()` method to analyze query execution plans and identify bottlenecks.
- **Aggregation Pipeline:** Optimize the aggregation pipeline stages to minimize data processing and improve efficiency.

19. Describe the Map-Reduce Functionality in MongoDB

Map-Reduce in **MongoDB** is a data processing paradigm used to perform complex data aggregation operations. It consists of two phases: the map phase processes each input document and emits key-value pairs, and the reduce phase processes all emitted values for each key and outputs the final result.

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).


```
db.collection.mapReduce(  
  function() { emit(this.category, this.price); },  
  function(key, values) { return Array.sum(values); },  
  { out: "category_totals" }  
);
```

This example calculates the total price for each category in a collection.

20. What is the Role of Journaling in MongoDB, and How Does It Impact Performance?

Journaling in MongoDB ensures data durability and crash recovery by recording changes to the data in a journal file before applying them to the database files. This mechanism allows MongoDB to recover from unexpected shutdowns or crashes by replaying the journal. While journaling provides data safety, it can impact performance due to the additional I/O operations required to write to the journal file.

21. How to Implement Full-Text Search in MongoDB?

Full-Text Search in MongoDB is implemented using text indexes. These indexes allow you to perform text search queries on string content within documents.

Example:

```
db.collection.createIndex({ content: "text" });  
db.collection.find({ $text: { $search: "mongodb" } });
```

In this example, a text index is created on the content field, and a text search query is performed to find documents containing the word "mongodb."

Considerations for deploying MongoDB in a production environment include:

- **Replication:** Set up replica sets for high availability and data redundancy.
- **Sharding:** Implement sharding for horizontal scaling and to distribute the load.
- **Backup and Recovery:** Establish a robust backup and recovery strategy.
- **Security:** Implement authentication, authorization, and encryption.
- **Monitoring:** Use monitoring tools to track performance and detect issues.
- **Capacity Planning:** Plan for adequate storage, memory, and CPU resources.
- **Maintenance:** Regularly update MongoDB to the latest stable version and perform routine maintenance tasks.

23. Explain the Concept of Horizontal Scalability and Its Implementation in MongoDB

- **Horizontal Scalability** in MongoDB refers to the ability to add more servers to distribute the load and data. This is achieved through sharding, where data is partitioned across multiple shards.
- Each shard is a replica set that holds a subset of the data. Sharding allows MongoDB to handle large datasets and high-throughput operations by distributing the workload.

24. How to Monitor and Troubleshoot Performance Issues in MongoDB?

Monitoring and troubleshooting performance issues in MongoDB involve:

- **Monitoring Tools:** Use tools like MongoDB Cloud Manager, MongoDB

- **Profiling:** Enable database profiling to capture detailed information about operations.
- **Explain Plans:** Use the `explain()` method to understand query execution and identify bottlenecks.
- **Index Analysis:** Review and optimize indexes based on query patterns and usage.
- **Resource Utilization:** Monitor CPU, memory, and disk I/O usage to identify resource constraints.

25. Describe the Process of Migrating Data from a Relational Database to MongoDB

Migrating data from a relational database to MongoDB involves several steps:

- **Schema Design:** Redesign the relational schema to fit MongoDB's document-oriented model. Decide on embedding vs. referencing, and plan for indexes and collections.
- **Data Export:** Export data from the relational database in a format suitable for MongoDB (e.g., CSV, JSON).
- **Data Transformation:** Transform the data to match the MongoDB schema. This can involve converting data types, restructuring documents, and handling relationships.
- **Data Import:** Import the transformed data into MongoDB using tools like `mongoimport` or custom scripts.
- **Validation:** Validate the imported data to ensure consistency and completeness.
- **Application Changes:** Update the application code to interact with MongoDB instead of the relational database.
- **Testing:** Thoroughly test the application and the database to ensure everything works as expected.

• **Go Live:** Deploy the MongoDB database in production and monitor the

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

MongoDB Query Based Interview Questions

MongoDB Query-Based Interview Questions focus on your ability to write efficient and optimized queries to interact with databases. These tasks include retrieving specific data using **filters**, **sorting** and **paginating results**, and **utilizing projections** to select desired fields. Below is a sample dataset we'll use to demonstrate various queries.

Sample Dataset

The following dataset represents a collection named **employees**, containing documents about employees in an organization. Each document includes details such as the **employee's name**, **age**, **position**, **salary**, **department**, and hire date.

```
"[
  {
    "_id": 1,
    "name": "John Doe",
    "age": 28,
    "position": "Software Engineer",
    "salary": 80000,
    "department": "Engineering",
    "hire_date": ISODate("2021-01-15")
  },
  {
    "_id": 2,
    "name": "Jane Smith",
    "age": 34,
    "position": "Project Manager",
    "salary": 95000,
    "department": "Engineering",
    "hire_date": ISODate("2019-06-23")
  },
  -
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
    "age": 41,
    "position": "CTO",
    "salary": 150000,
    "department": "Management",
    "hire_date": ISODate("2015-03-12")
  },
  {
    "_id": 4,
    "name": "Michael Brown",
    "age": 29,
    "position": "Software Engineer",
    "salary": 85000,
    "department": "Engineering",
    "hire_date": ISODate("2020-07-30")
  },
  {
    "_id": 5,
    "name": "Sarah Davis",
    "age": 26,
    "position": "UI/UX Designer",
    "salary": 70000,
    "department": "Design",
    "hire_date": ISODate("2022-10-12")
  }
]
```

1. Find all Employees Who Work in the "Engineering" Department.

Query:

```
db.employees.find({ department: "Engineering" })
```

Output:

```
[
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#)

```
"age": 28,  
"position": "Software Engineer",  
"salary": 80000,  
"department": "Engineering",  
"hire_date": ISODate("2021-01-15")  
},  
{  
  "_id": 2,  
  "name": "Jane Smith",  
  "age": 34,  
  "position": "Project Manager",  
  "salary": 95000,  
  "department": "Engineering",  
  "hire_date": ISODate("2019-06-23")  
},  
{  
  "_id": 4,  
  "name": "Michael Brown",  
  "age": 29,  
  "position": "Software Engineer",  
  "salary": 85000,  
  "department": "Engineering",  
  "hire_date": ISODate("2020-07-30")  
}  
]
```

Explanation: This query finds all employees whose department field is "Engineering".

2. Find the Employee with the Highest Salary.

Query:

```
db.employees.find().sort({ salary: -1 }).limit(1)
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
[
  {
    "_id": 3,
    "name": "Emily Johnson",
    "age": 41,
    "position": "CTO",
    "salary": 150000,
    "department": "Management",
    "hire_date": ISODate("2015-03-12")
  }
]
```

Explanation: This query sorts all employees by salary in descending order and retrieves the top document, which is the employee with the highest salary.

3. Update the Salary of "John Doe" to 90000.

Query:

```
db.employees.updateOne({ name: "John Doe" }, { $set: { salary:
90000 } })
```

Output:

```
{ "acknowledged" : true, "matchedCount" : 1, "modifiedCount" :
1 }
```

Explanation: This query updates the salary of the employee named "John Doe" to 90000.

4. Count the Number of Employees in Each Department.

Query:

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Output:

Explanation: This query groups the employees by the department field and counts the number of employees in each department.

5. Add a New Field Bonus to All Employees in the "Engineering" Department with a Value of 5000.

```
db.employees.updateMany({ department: "Engineering" }, { $set:
{ bonus: 5000 } })
```

```
{ "acknowledged" : true, "matchedCount" : 3, "modifiedCount" : 3 }
```

Explanation: This query adds a new field bonus with a value of 5000 to all employees in the "Engineering" department.

6. Retrieve All Documents in the Employees Collection and Sort Them by the Length of Their Name in Descending Order.

11 7 1 / F

<https://www.geeksforgeeks.org/mongodb/mongodb-interview-questions/>


```
{ $project: { nameLength: 0 } }  
])
```

Output:

```
[  
  {  
    "_id": 2,  
    "name": "Jane Smith",  
    "age": 34,  
    "position": "Project Manager",  
    "salary": 95000,  
    "department": "Engineering",  
    "hire_date": ISODate("2019-06-23")  
  },  
  {  
    "_id": 3,  
    "name": "Emily Johnson",  
    "age": 41,  
    "position": "CTO",  
    "salary": 150000,  
    "department": "Management",  
    "hire_date": ISODate("2015-03-12")  
  },  
  {  
    "_id": 1,  
    "name": "John Doe",  
    "age": 28,  
    "position": "Software Engineer",  
    "salary": 80000,  
    "department": "Engineering",  
    "hire_date": ISODate("2021-01-15")  
  },  
  {  
    "_id": 4,  
    "name": "Michael Brown"
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
"department": "Engineering",
"hire_date": ISODate("2020-07-30")
},
{
  "_id": 5,
  "name": "Sarah Davis",
  "age": 26,
  "position": "UI/UX Designer",
  "salary": 70000,
  "department": "Design",
  "hire_date": ISODate("2022-10-12")
}
]
```

Explanation: This query calculates the length of each employee's name, sorts the documents by this length in descending order, and removes the temporary nameLength field from the output.

7. Find the Average Salary of Employees in the "Engineering" Department.

Query:

```
db.employees.aggregate([
  { $match: { department: "Engineering" } },
  { $group: { _id: null, averageSalary: { $avg: "$salary" } } }
])
```

Output:

```
[
  { "_id": null, "averageSalary": 86666.66666666667 }
]
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

8. Find the Department with the Highest Average Salary.

Query:

```
db.employees.aggregate([
  { $group: { _id: "$department", averageSalary: { $avg:
"$salary" } } },
  { $sort: { averageSalary: -1 } },
  { $limit: 1 }
])
```

Output:

```
[
  { "_id": "Management", "averageSalary": 150000 }
]
```

Explanation: This query groups employees by department, calculates the average salary for each department, sorts these averages in descending order, and retrieves the department with the highest average salary.

9. Find the Total Number of Employees Hired in Each Year.

Query:

```
db.employees.aggregate([
  { $group: { _id: { $year: "$hire_date" }, totalHired: {
$sum: 1 } } }
])
```

Output:

```
[
  { "_id": 2015, "totalHired": 1 },
  { "_id": 2019, "totalHired": 1 },
]
```

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

```
{ "_id": 2022, "totalHired": 1 }  
]
```

Explanation: This query groups employees by the year they were hired, which is extracted from the hire_date field, and counts the total number of employees hired each year.

10. Find the Highest and Lowest Salary in the "Engineering" Department.

Query:

```
db.employees.aggregate([  
  { $match: { department: "Engineering" } },  
  {  
    $group: {  
      _id: null,  
      highestSalary: { $max: "$salary" },  
      lowestSalary: { $min: "$salary" }  
    }  
  }  
])
```

Output:

```
[  
  { "_id": null, "highestSalary": 95000, "lowestSalary":  
    80000 }  
]
```

Explanation: This query filters employees to those in the "Engineering" department, then calculates the highest and lowest salary within this group.

Conclusion

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

interview questions, covering **fundamental concepts**, **advanced topics**, and **real-world applications**. By practicing these questions, you'll enhance your knowledge of MongoDB, strengthen your **problem-solving skills**, and build the confidence needed to excel in your interview. Stay consistent in your preparation, and you'll be ready to tackle any **MongoDB challenge** that comes your way!

[Comment](#)[More info](#)[Advertise with us](#)

Next Article

Top 50 MongoDB Interview
Questions with Answers for 2025

Similar Reads

1. [Top 50 Database Interview Questions and Answers for 2025](#)
2. [Top 50 PostgreSQL Interview Questions and Answers](#)
3. [Top 50 SQLite Interview Questions for 2024](#)
4. [30 SQL Interview Questions For Business Analyst in 2025](#)
5. [Top 10 MongoDB Tools for 2025](#)
6. [How to Manage Data with MongoDB](#)
7. [MongoDB Roadmap: A Complete Guide \[2025 Updated\]](#)
8. [A Better Developer Experience with the MongoDB Atlas Data API](#)
9. [How to Migrate a Microservice from MySQL to MongoDB](#)
10. [What is the Format of Document in MongoDB?](#)



We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Registered Address:

K 061, Tower K, Gulshan Vivante
Apartment, Sector 137, Noida, Gautam
Buddh Nagar, Uttar Pradesh, 201305



Advertise with us

Company

About Us
Legal
Privacy Policy
Careers
In Media
Contact Us
Corporate Solution
Campus Training Program

Explore

Job-A-Thon
Offline Classroom Program
DSA in JAVA/C++
Master System Design
Master CP
Videos

Tutorials

Python
Java
C++
PHP
GoLang
SQL
R Language
Android

DSA

Data Structures
Algorithms
DSA for Beginners
Basic DSA Problems
DSA Roadmap
DSA Interview Questions
Competitive Programming

Data Science & ML

Data Science With Python
Machine Learning
ML Maths
Data Visualisation
Pandas
NumPy
NLP
Deep Learning

Web Technologies

HTML
CSS
JavaScript
TypeScript
ReactJS
NextJS
NodeJs
Bootstrap

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

Django Tutorial
Python Projects
Python Tkinter
Web Scraping
OpenCV Tutorial
Python Interview Question

DevOps

Git
AWS
Docker
Kubernetes
Azure
GCP
DevOps Roadmap

School Subjects

Mathematics
Physics
Chemistry
Biology
Social Science
English Grammar

Preparation Corner

Company-Wise Recruitment Process
Aptitude Preparation
Puzzles
Company-Wise Preparation

Courses

IBM Certification Courses
DSA and Placements
Web Development
Data Science
Programming Languages
DevOps & Cloud

Operating Systems
Computer Network
Database Management System
Software Engineering
Digital Logic Design
Engineering Maths

System Design

High Level Design
Low Level Design
UML Diagrams
Interview Guide
Design Patterns
OOAD
System Design Bootcamp
Interview Questions

Databases

SQL
MYSQL
PostgreSQL
PL/SQL
MongoDB

More Tutorials

Software Development
Software Testing
Product Management
Project Management
Linux
Excel
All Cheat Sheets

Programming Languages

C Programming with Data Structures
C++ Programming Course
Java Programming Course
Python Full Course

Clouds/Devops

GATE 2026

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).

@GeeksforGeeks, Sanchhaya Education Private Limited, All rights reserved

We use cookies to ensure you have the best browsing experience on our website. By using our site, you acknowledge that you have read and understood our [Cookie Policy](#) & [Privacy Policy](#).